

Build a Table-centric Apache Flink Ecosystem

Shaoxuan Wang
Director, Senior Staff Engineer at Alibaba

2019 Flink Forward San Francisco





Shaoxuan Wang
Director, Senior Staff Engineer at
Alibaba

shaoxuan.wsx@alibaba-inc.com
shaoxuan@apache.org

Peking University

EECS

**University of California
at San Diego**

PhD, Computer Engineer

Broadcom

High-Perf Platform
Senior Software Engineer

Facebook

Core Infra
Software Engineer

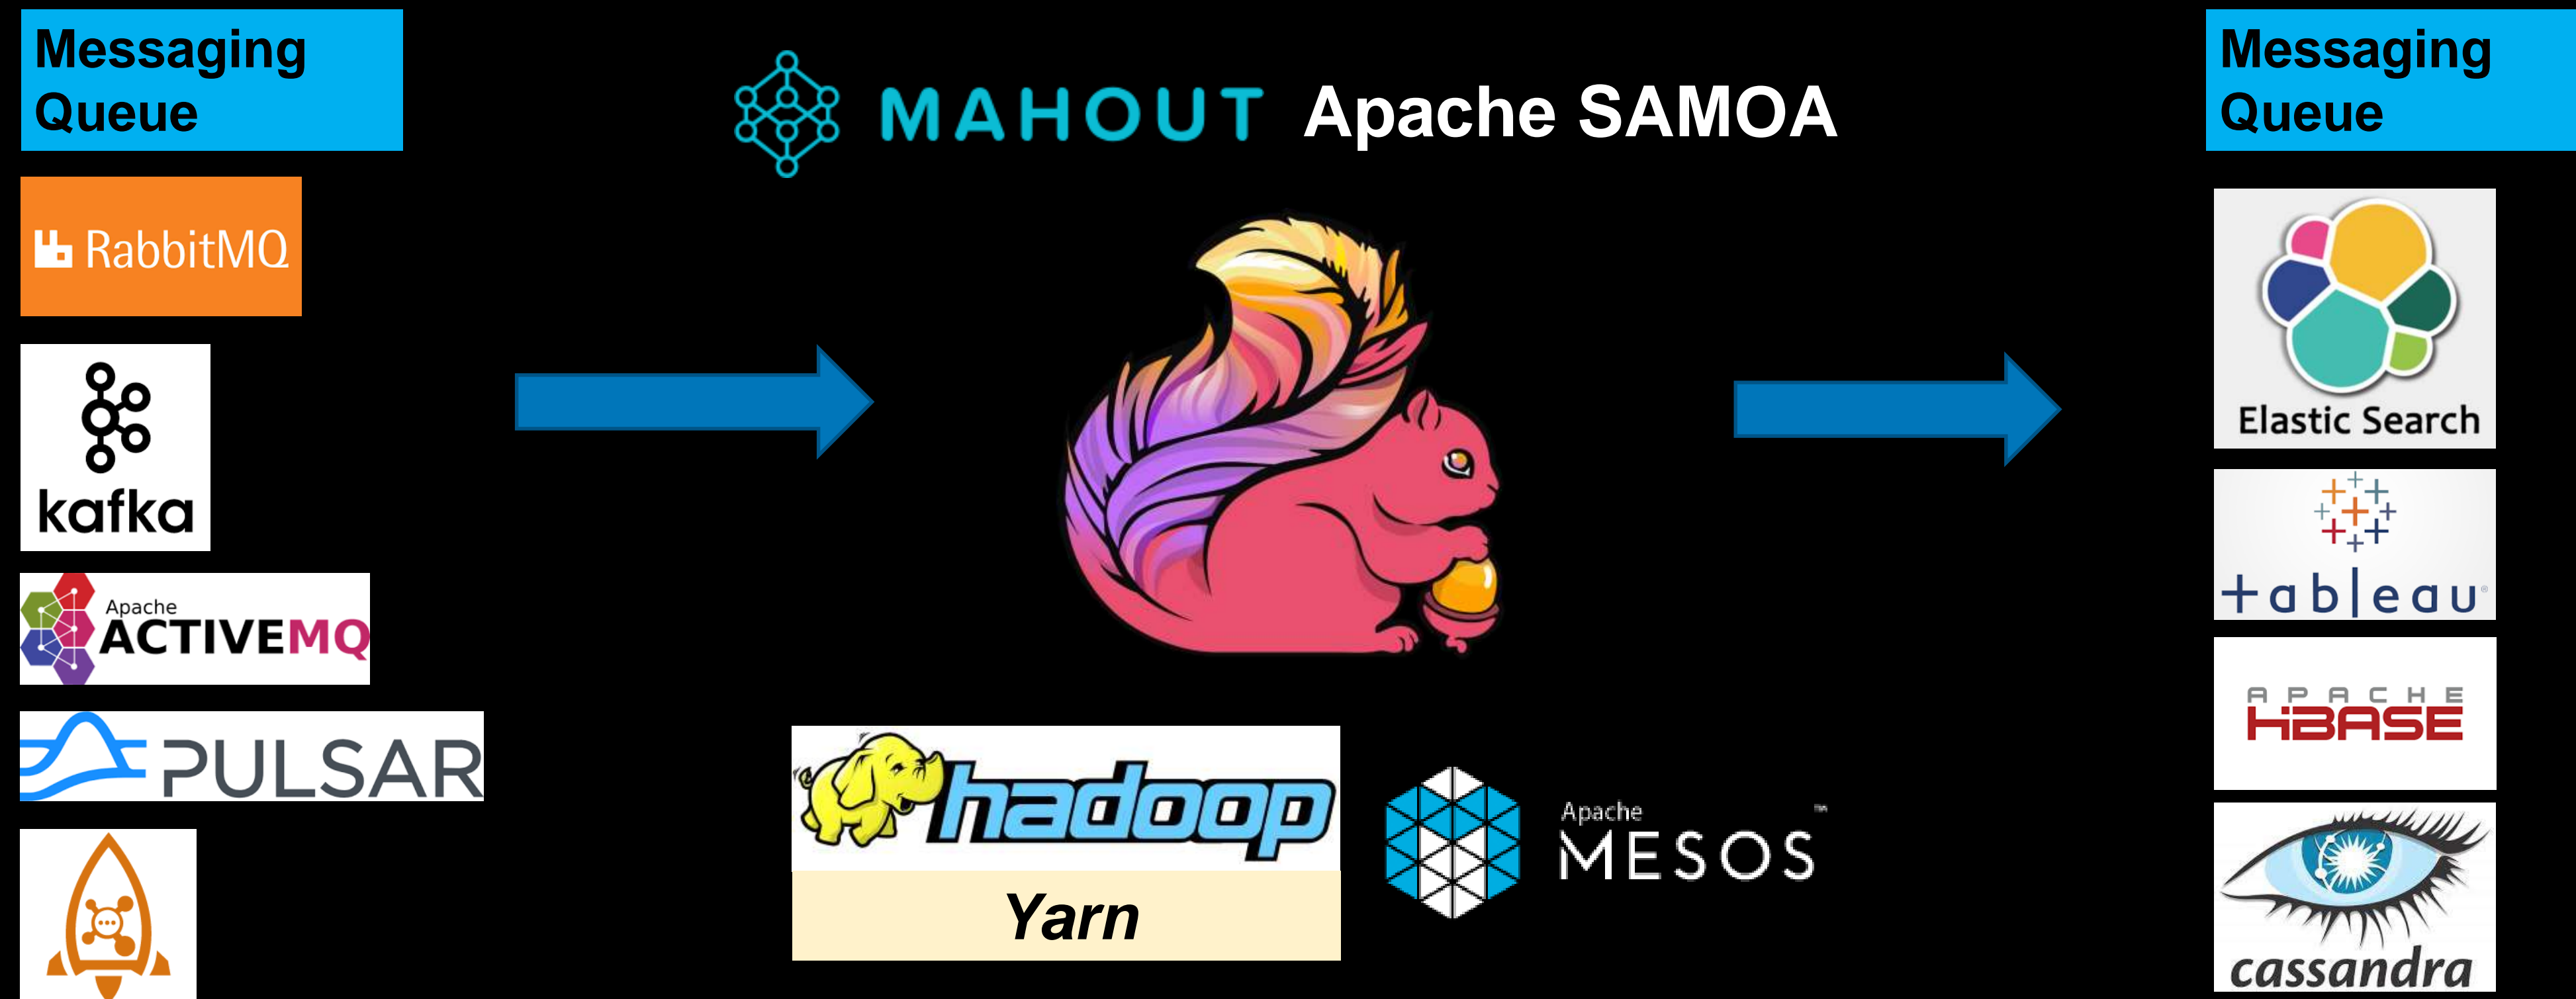
Alibaba Group

Data Infra / Cloud
Senior Staff Engineer

Flink Committer

Since 2017

Apache Flink Ecosystem - Present



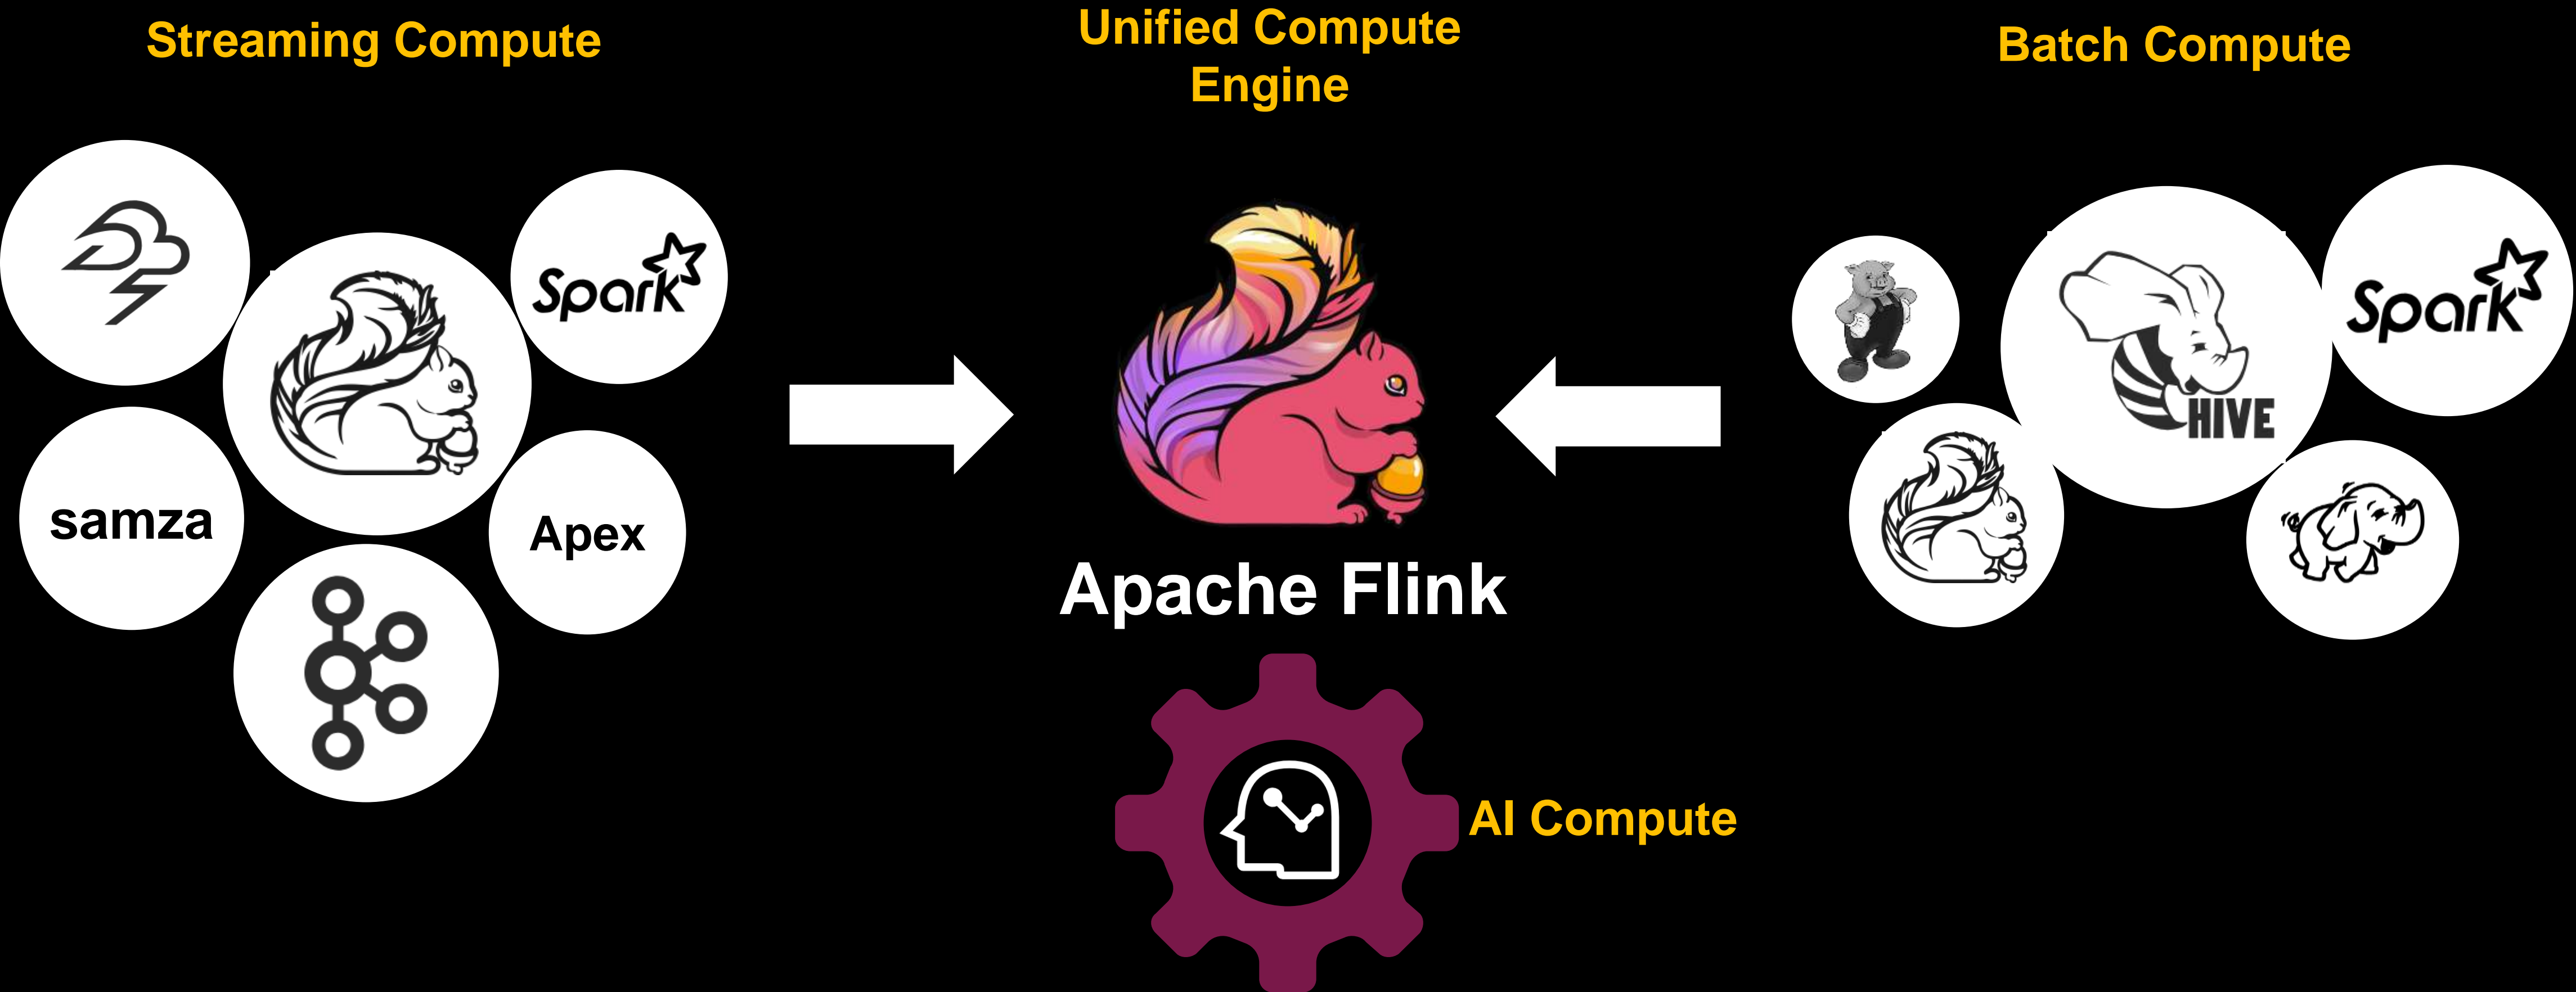
Apache Flink is the most sophisticated open-source Stream Processor

Intelligent Big Data Computing

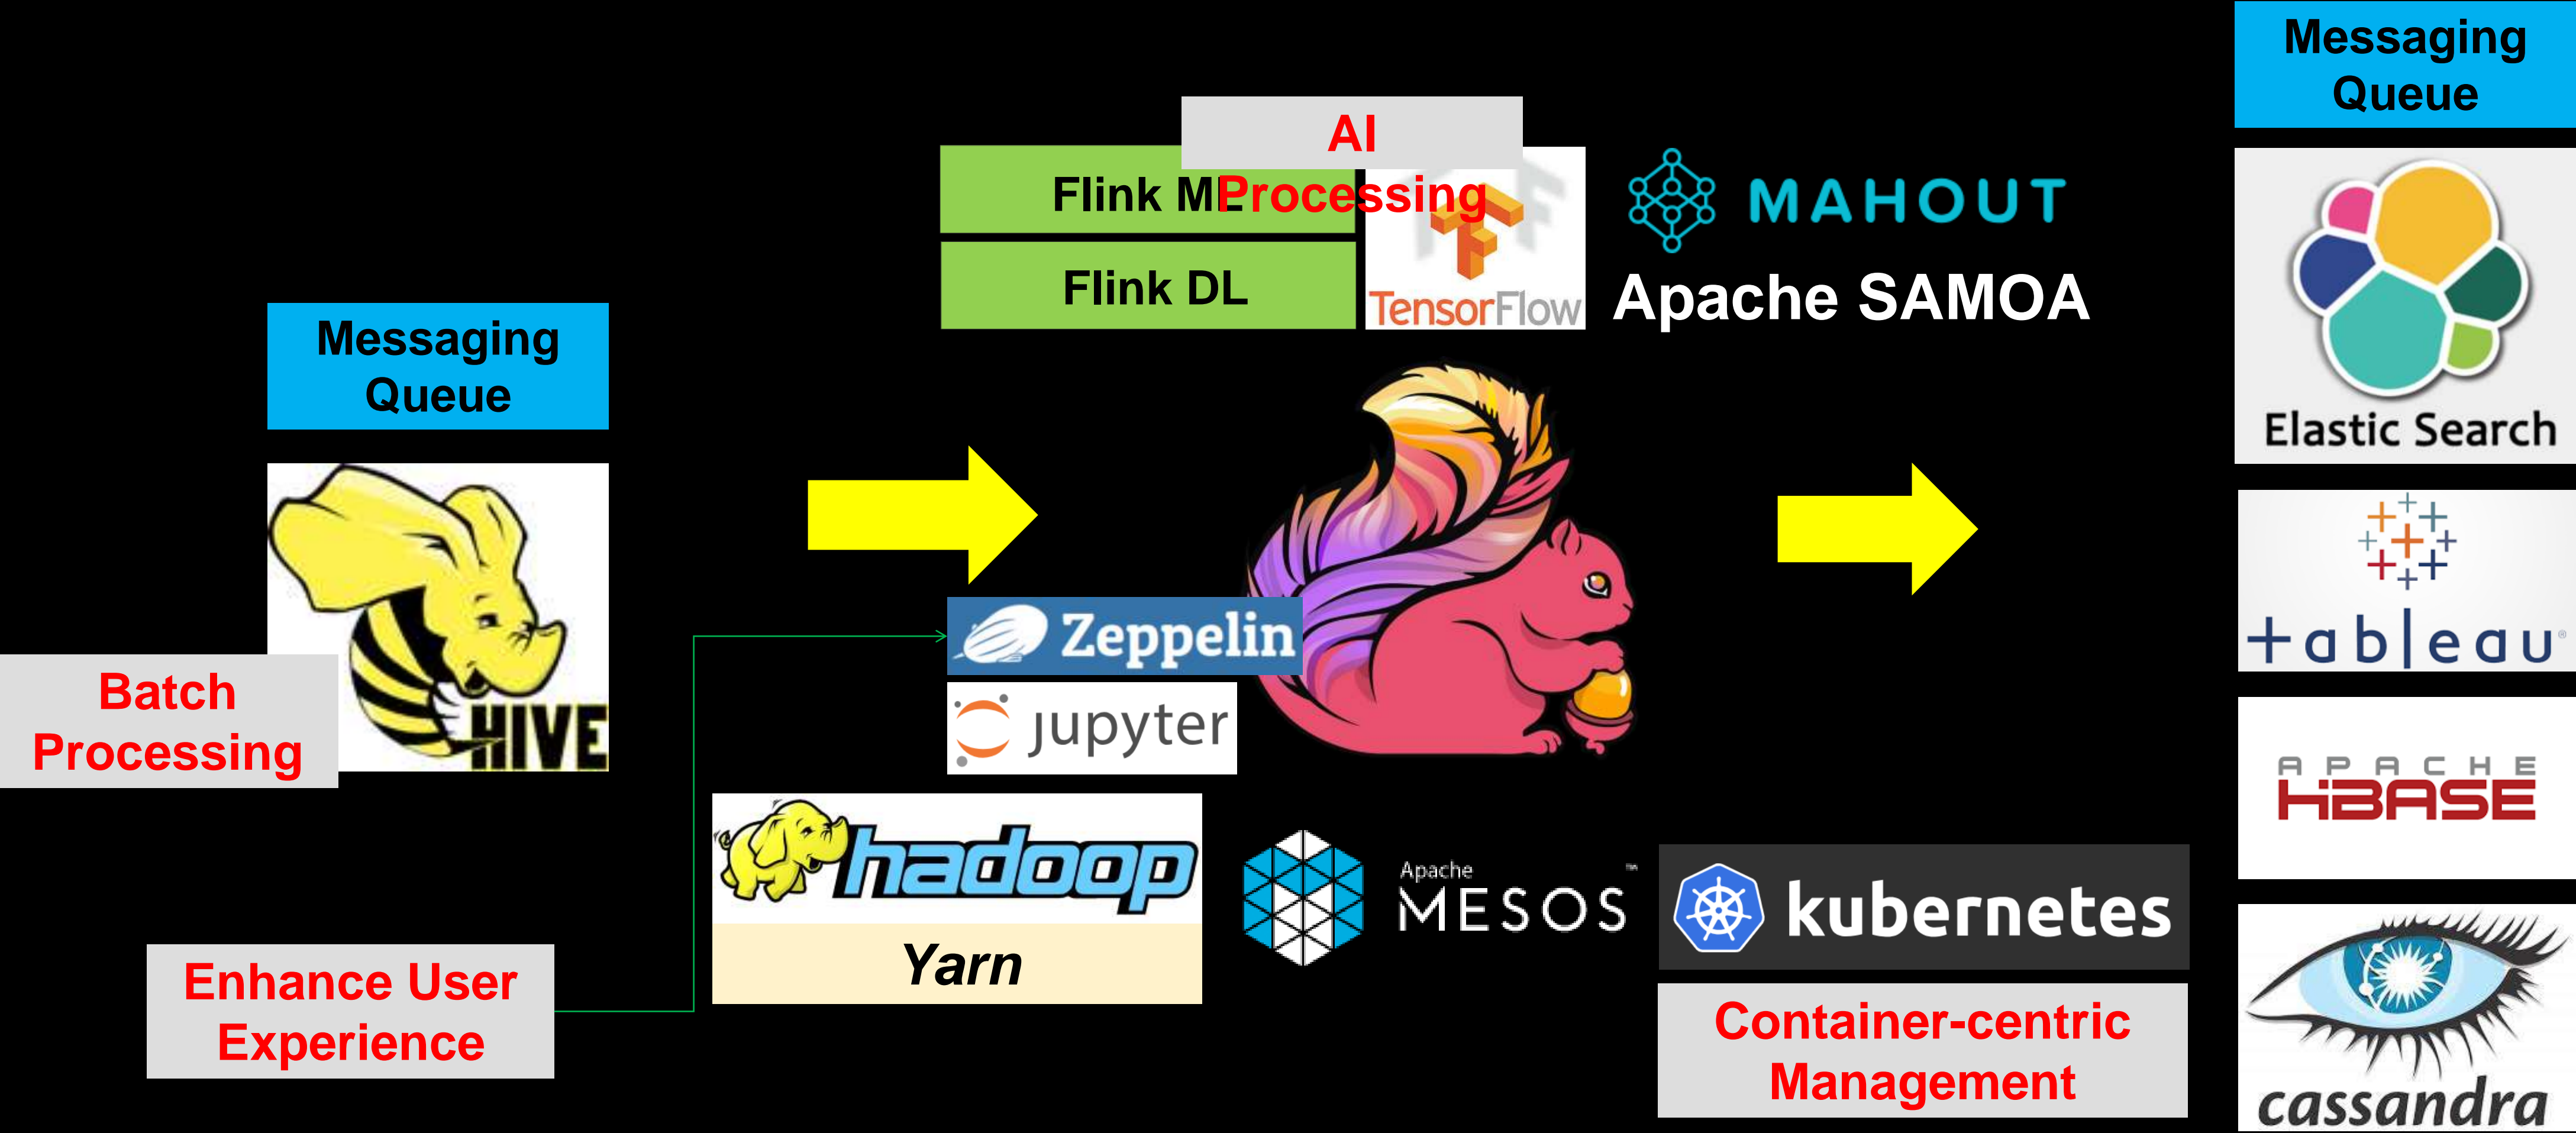


Can Apache Flink become an unified engine for intelligent big data computing?

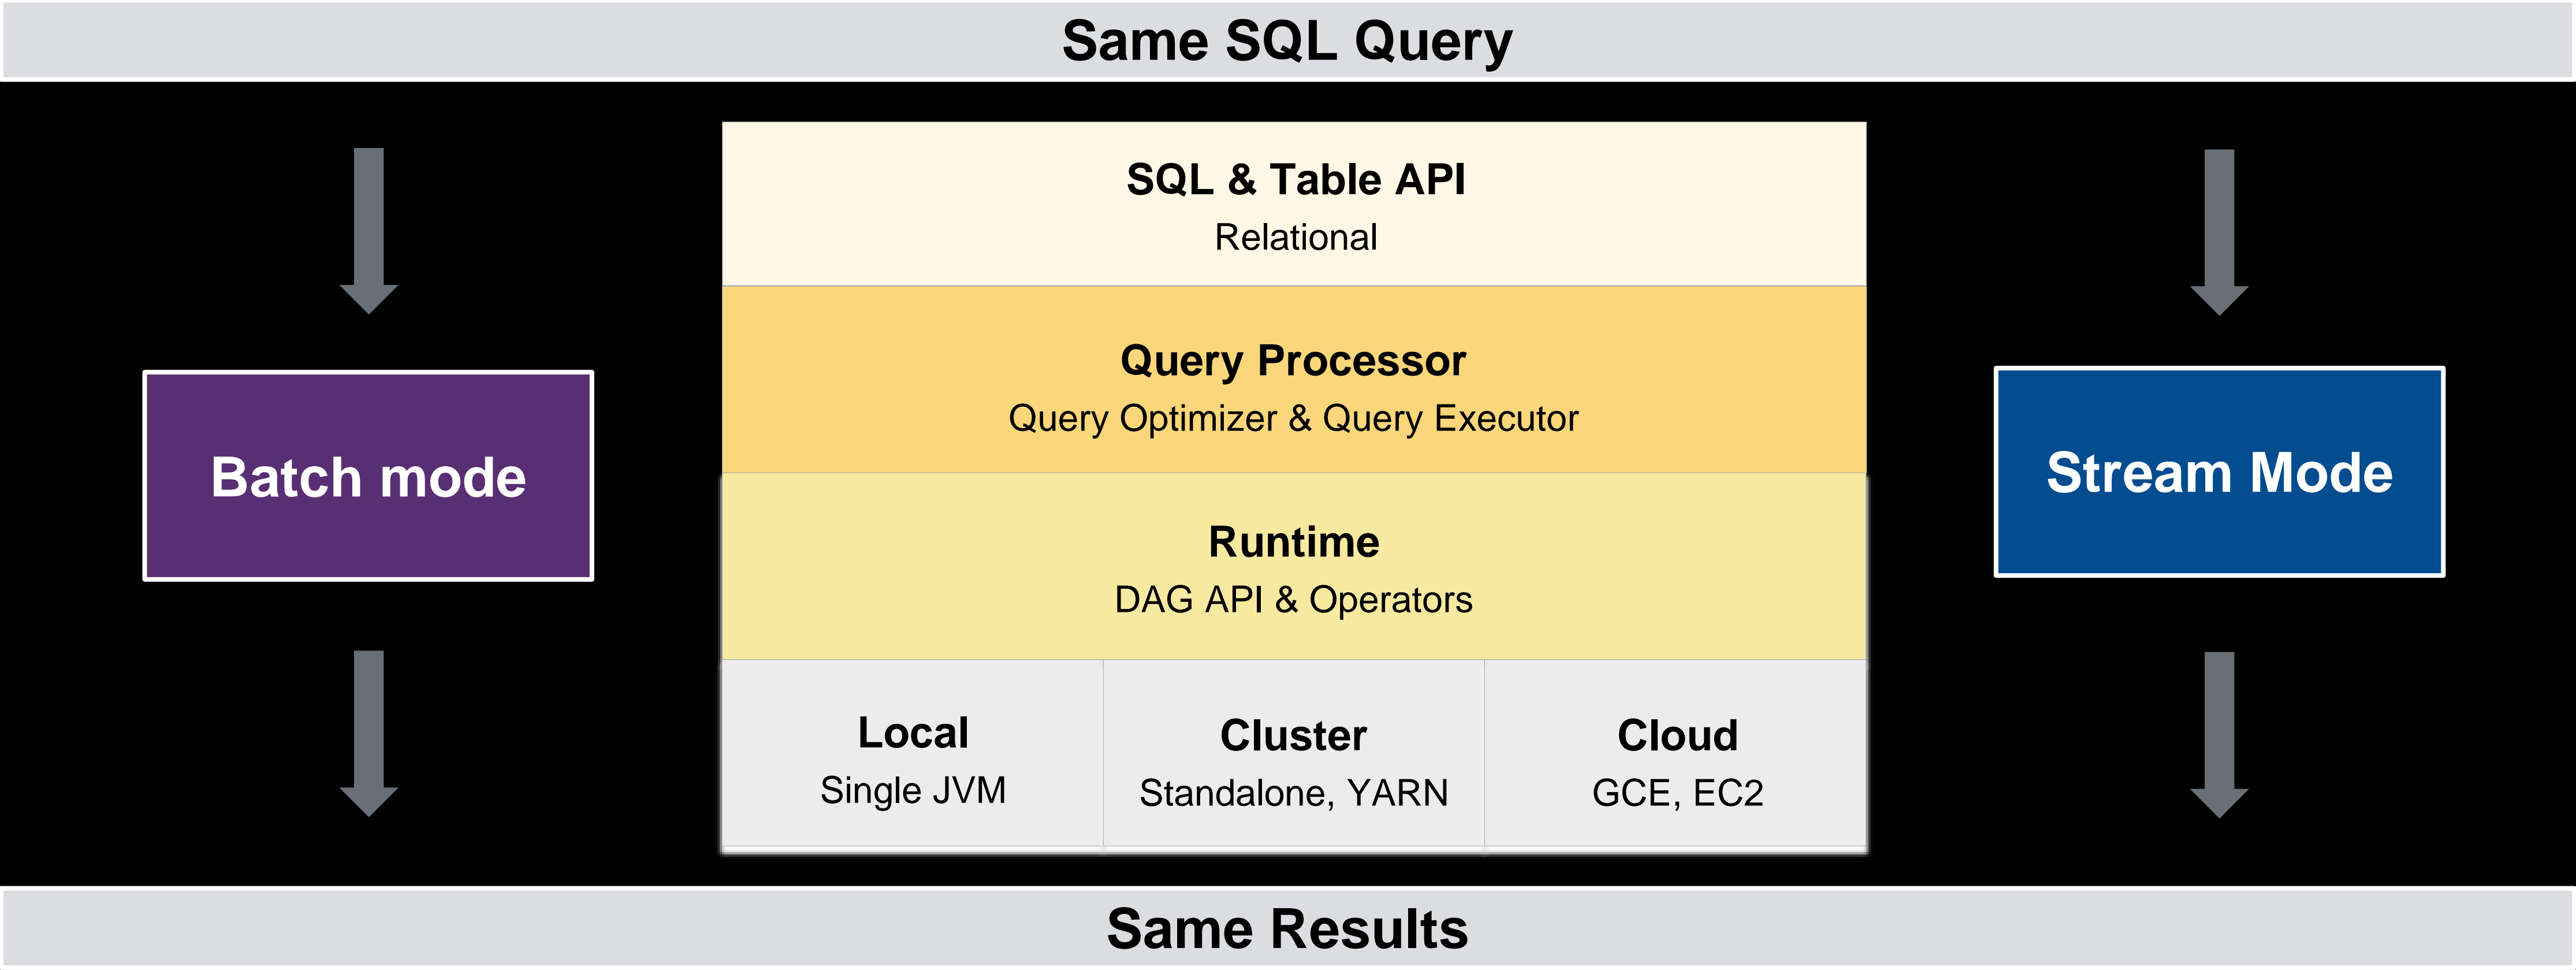
Build Intelligent Big Data Platform with Apache Flink



Apache Flink Ecosystem - Future



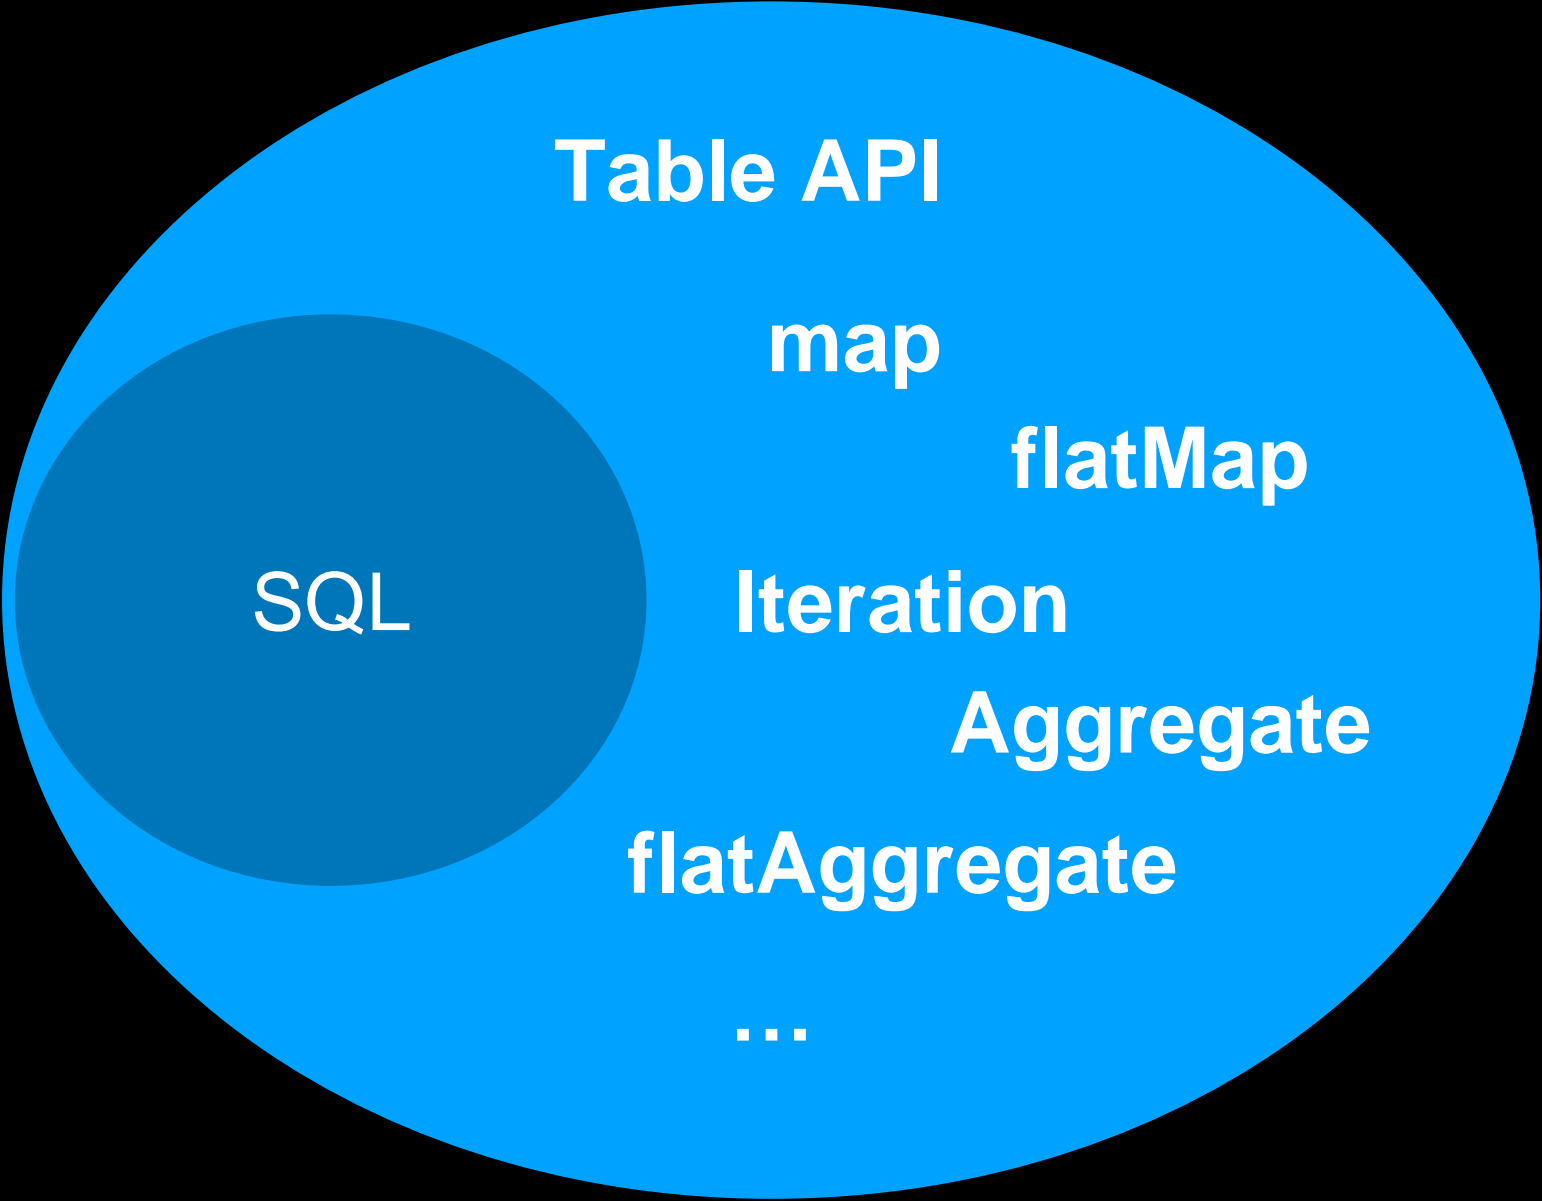
TableAPI: declarative API with nature Optimization



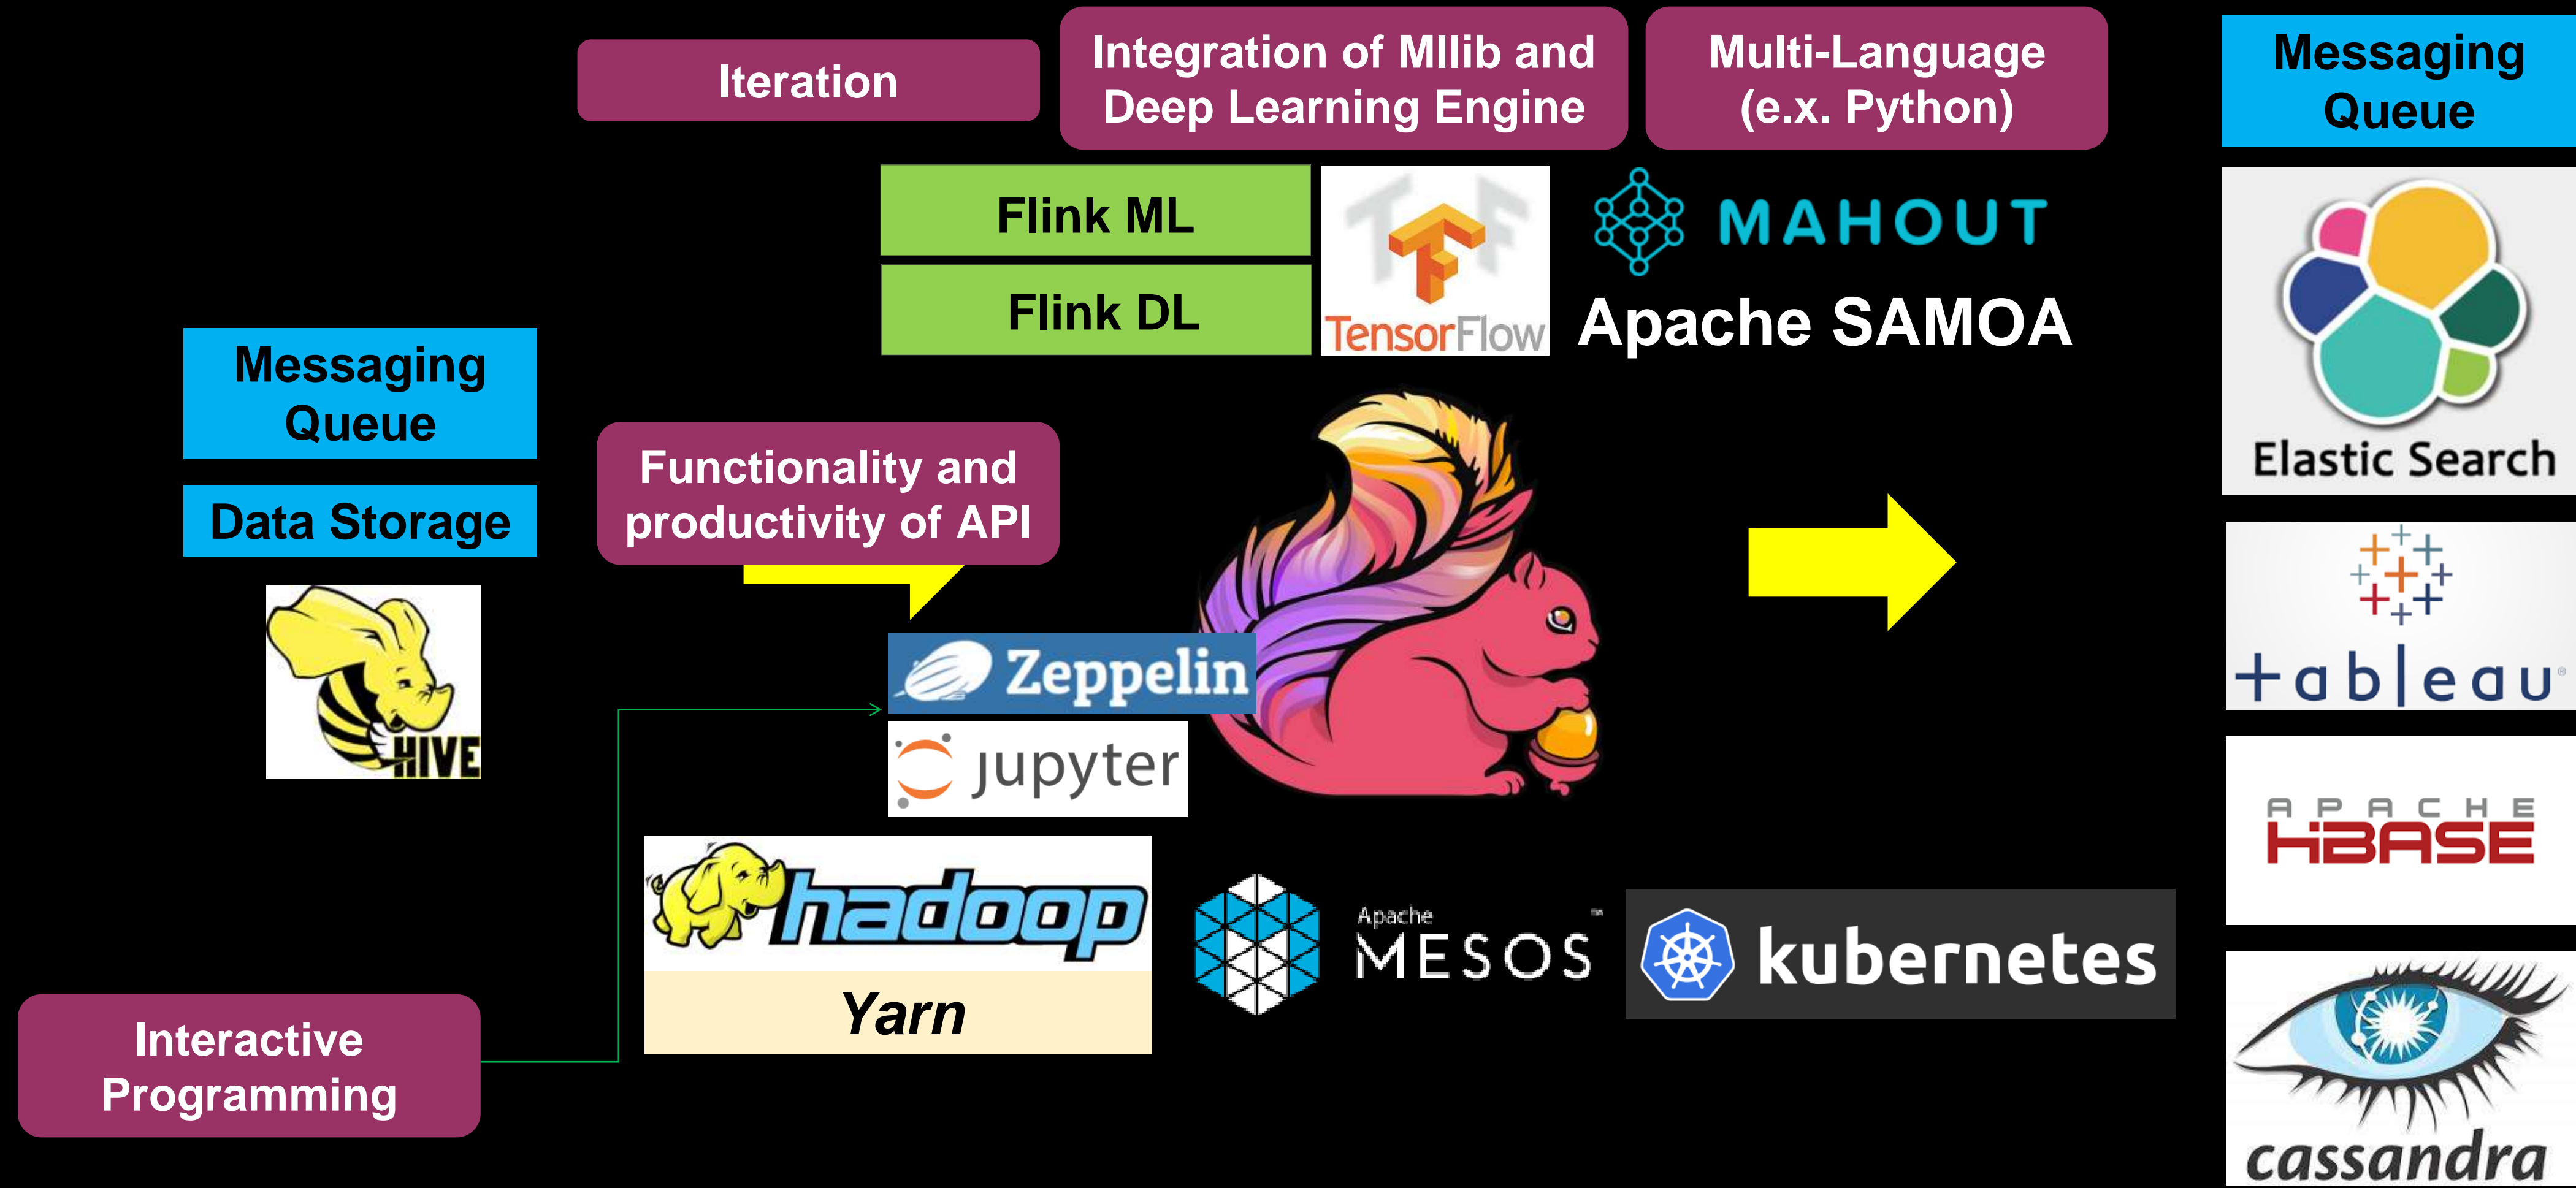
similar as SQL, batch&stream unified, declarative API with nature optimization framework

TableAPI is more than SQL

| | TableAPI | SQL | e.g. |
|--------------------------|----------|-----|--------------------------------|
| Stream and batch unified | Y | Y | SELECT/AGG/WIN DOW etc. |
| Functional scalability | Y | N | flatAgg/Column operations etc. |
| Rich expression | Y | N | map/flatMap/intersect etc. |
| Compile check | Y | N | Java/Scala IDE |



What Else are Needed on TableAPI



What Else are Needed on TableAPI



**Functionality
and Productivity
Enhancement**



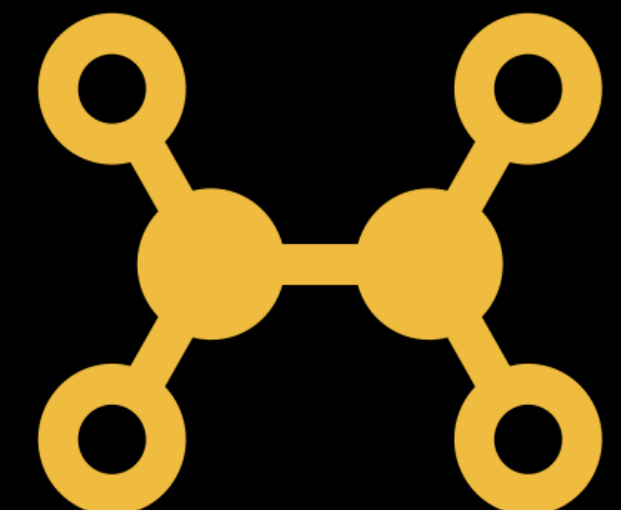
**Interactive
programming**



Multi-language



Iteration

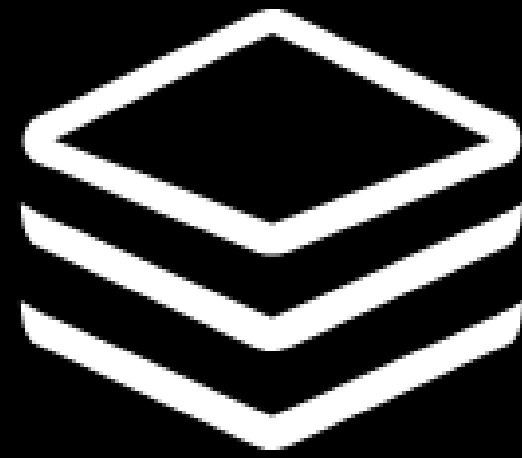


**Execute MLlib
and DL engine**

Agenda

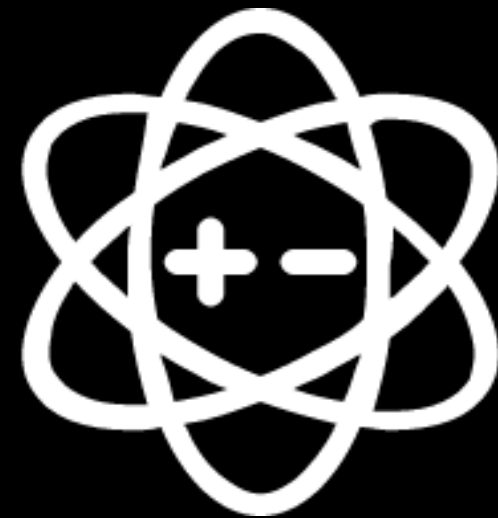
- **TableAPI enhancement: functionality and productivity**
- **Interactive programming**
- **Multi-language support**
- **Brief introduction for AI support**

TableAPI Enhancement



**Row-based data
processing API**

**Functionality,
productivity**



**Column
operation**

Productivity

Hint

**optimization instruction,
Resource configuration**

Introduce Row-based TableAPI (FLIP29)

| UDF | Single Row Input | Multiple Row Input |
|---------------------|------------------|---|
| Single Row Output | ScalarFunction | AggregateFunction |
| Multiple Row Output | TableFunction | <i>TableAggregateFunction (new)</i> |

| Table Method | Single Column Output | Multiple Column Output |
|------------------------|----------------------|-----------------------------|
| ScalarFunction | Table.select | <i>Table.map</i> |
| AggregateFunction | Table.select | <i>GroupedTable.agg</i> |
| TableFunction | N/A | <i>Table.flatmap</i> |
| TableAggregateFunction | N/A | <i>GroupedTable.flatagg</i> |

Map Operator in Table API

Method signature

`def map(scalaFunction: Expression): Table`

Pros

`table.select(udf('c1), udf('c2), udf('c3))`

VS

`table.map(udf('c1,'c2,'c3))`

Example

```
val res = tab
  .map(fun('e)).as('a, 'b, 'c)
  .select('a, 'c)
```

```
class MyMap extends ScalarFunction {
  var param : String = ""

  override def open(context: FunctionContext): Unit
    = param = context.getJobParameter("paramKey","")

  def eval(index: Int): Row = {
    val result = new Row(3)
    // Business processing based on data and parameters
    result
  }

  override def getResultType(signature: Array[Class[_]]): TypeInformation[_]
    = {
      Types.ROW(Types.STRING, Types.INT, Types.LONG)
    }
  }
```

| INPUT(Row) | OUTPUT(Row) |
|------------|-------------|
| 1 | 1 |

FatMap Operator in Table API

Method signature

`def flatMap(tableFunction: Expression): Table`

Pros

`table.join(udtf)` VS `table.flatMap(udtf())`

Example

```
val res = tab
  .flatMap(fun('e,'f)).as('name, 'age)
  .select('name, 'age)
```

```
case class User(name: String, age: Int)

class MyFlatMap extends TableFunction[User] {
  def eval(user: String): Unit = {
    if (user.contains("#")) {
      val splits = user.split("#")
      collect(User(splits(0),splits(1).toInt))
    }
  }
}
```

| INPUT(Row) | OUTPUT(Row) |
|------------|-------------|
| 1 | N(N>=0) |

Aggregate Operator in Table API

Method signature

```
def aggregate(aggregateFunction: Expression): AggregatedTable
class AggregatedTable(table: Table, groupKeys: Seq[Expression], aggFunction: Expression)
```

Pros

```
table.select(agg1(), agg2(), agg3()....)
VS
table.aggregate(agg())
```

Example

```
val res = groupedTab
    .groupBy('a)
    .aggregate(
        aggFun('e,'f) as ('a, 'b, 'c))
    .select('a, 'c)
```

```
class CountAccumulator extends JTuple1[Long] {
    f0 = 0L //count
}
class CountAgg extends AggregateFunction[JLong, CountAccumulator] {
    def accumulate(acc: CountAccumulator): Unit = {
        acc.f0 += 1L
    }

    override def getValue(acc: CountAccumulator): JLong = {
        acc.f0
    }
    ... retract()/merge()
}
```

| INPUT(Row) | OUTPUT(Row) |
|------------|-------------|
| N(N>=0) | 1 |

FlatAggregate Operator in Table API

Method signature

```
def flatAggregate(tableAggregateFunction: Expression): GroupedFlatAggregateTable
class GroupedFlatAggTable(table: Table, groupKey: Seq[Expression], tableAggFun: Expression)
```

Pros

A completely new feature

Example

```
val res = groupedTab
  .groupBy('a)
  .faltAggregate(
    flatAggFun('e,'f) as ('a, 'b, 'c))
  .select('a, 'c)
```

```
class TopNAcc {
  var data: MapView[JInt, JLong] = _ // (rank -> value)
  ...
}

class TopN(n: Int) extends TableAggregateFunction[(Int, Long), TopNAccum] {

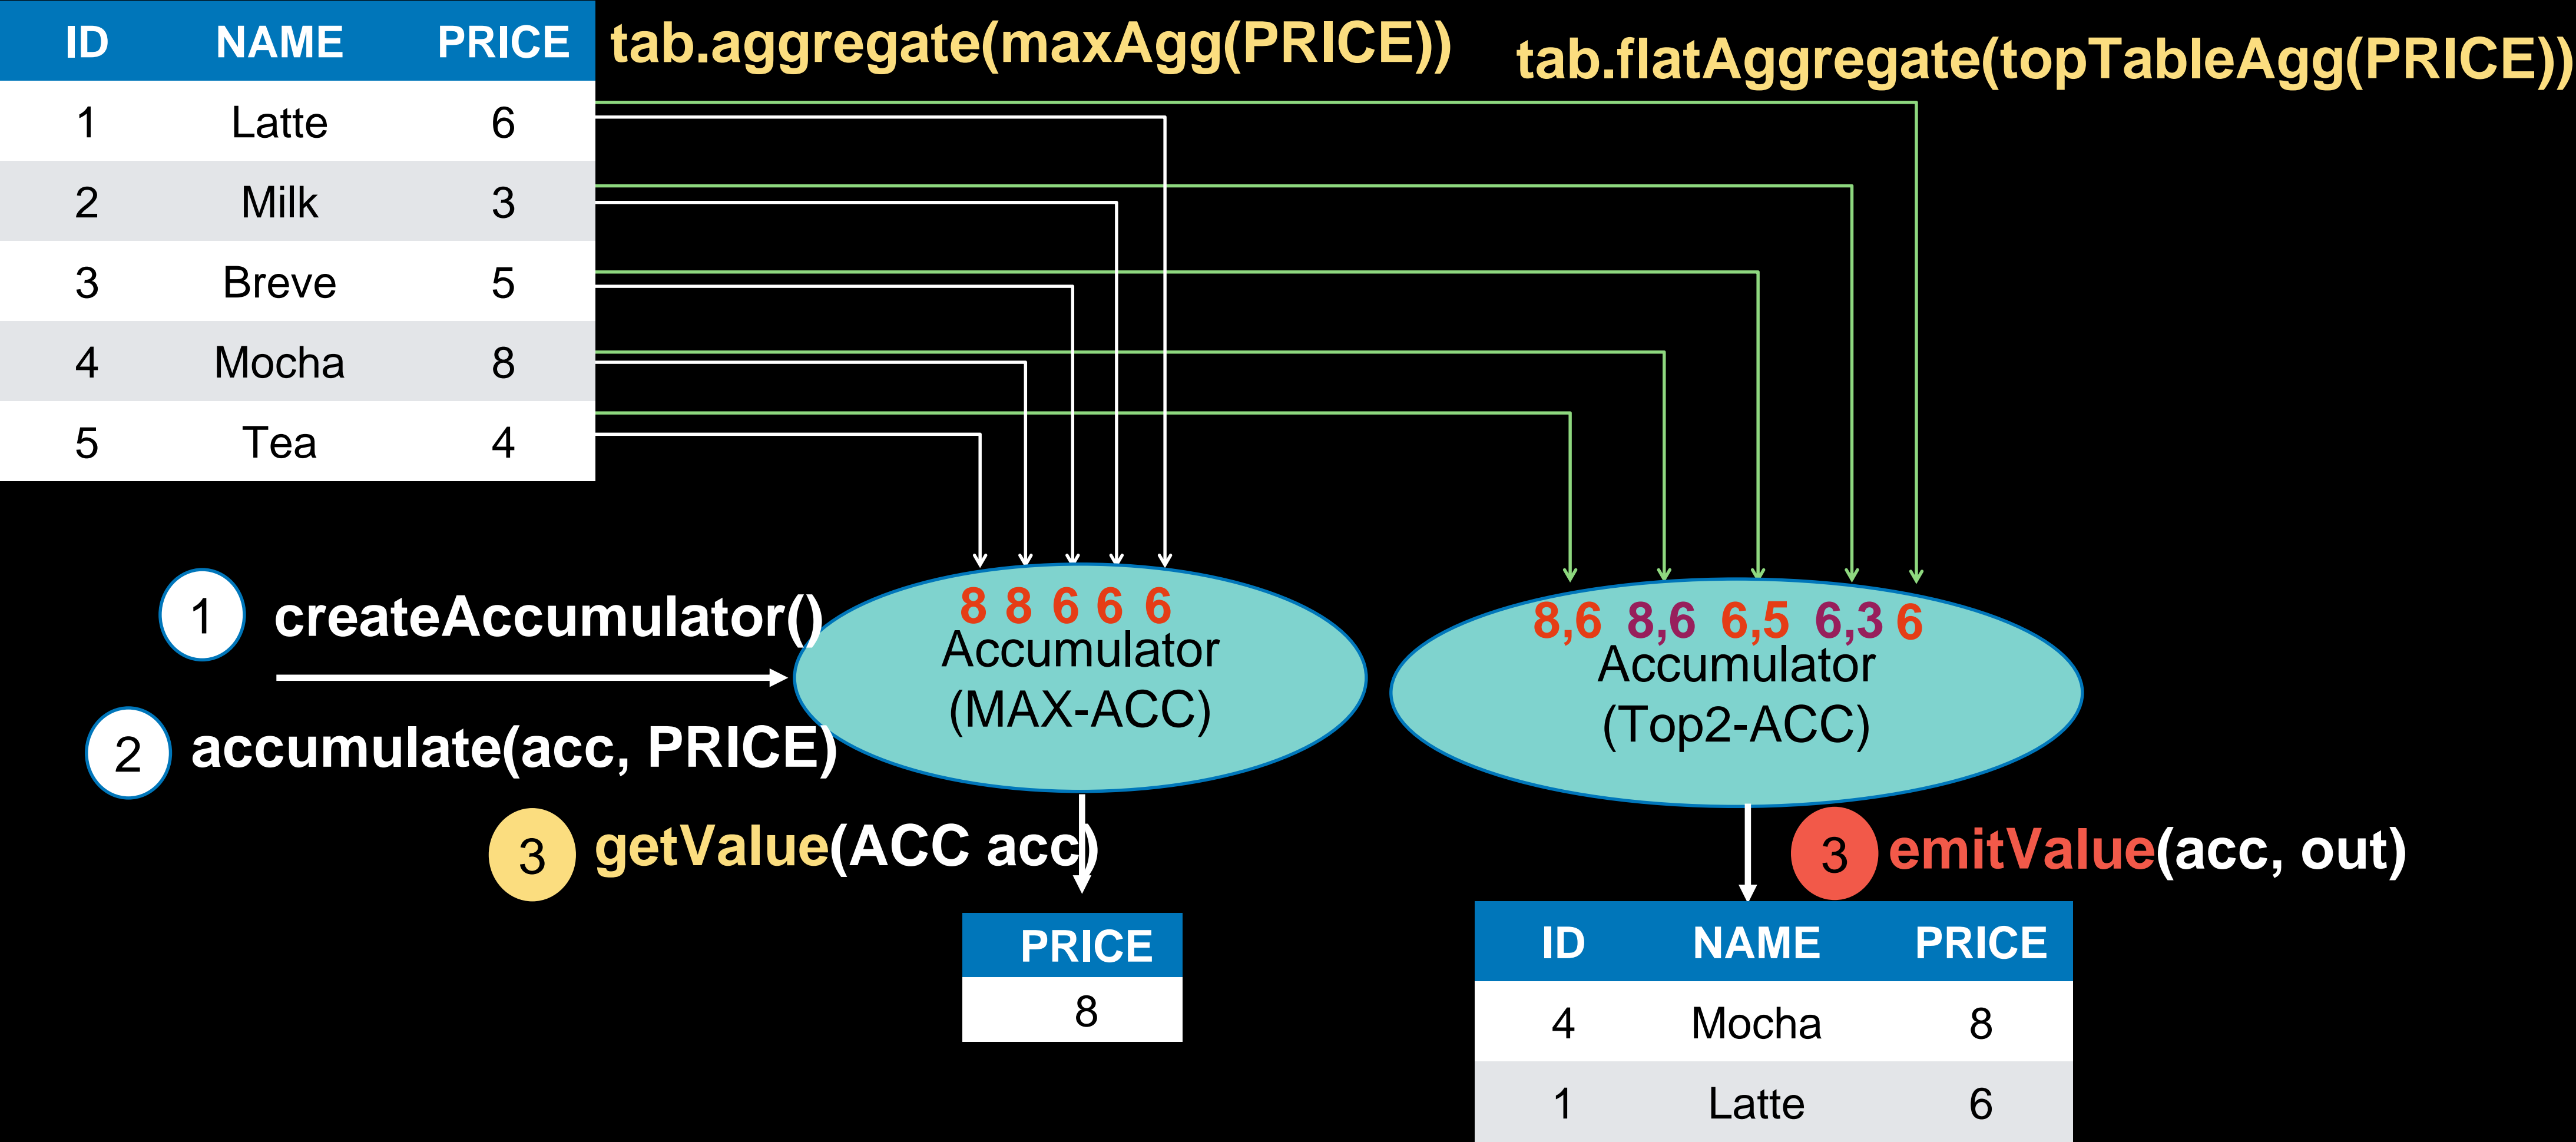
  def accumulate(acc: TopNAcc, category: Int, value: Long) {
    ...
  }

  def emitValue(acc: TopNAcc, out: Collector[(Int, Long)]): Unit = {
    ...
  }
  ...retract/merge
}
```

| INPUT(Row) | OUTPUT(Row) |
|------------|-------------|
| N(N>=0) | M(M>=0) |

Examples of Aggregate and FlatAggregate

Examples of Max (aggregate) and Top2 (flatAggregate)

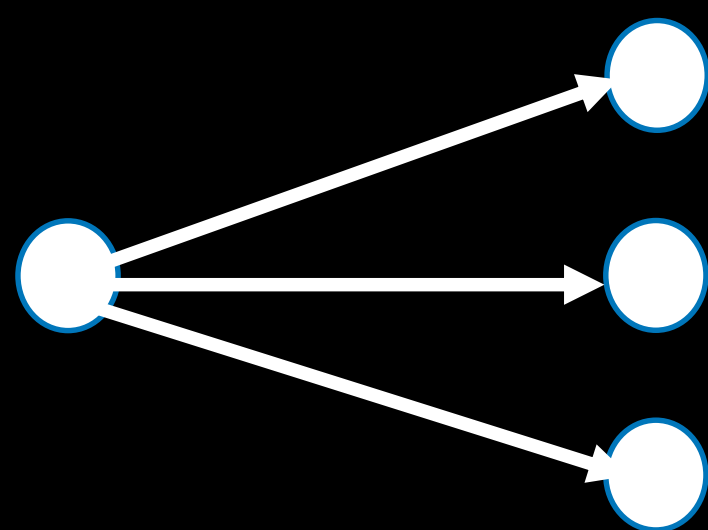


Summary

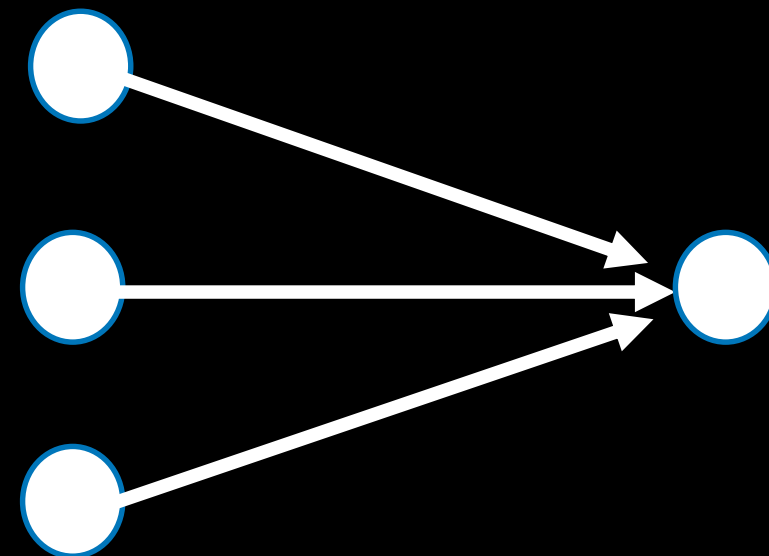
| | Single Row Input | Multiple Row Input |
|---------------------|---------------------------------------|---|
| Single Row Output | ScalarFunction (select/map) | AggregateFunction (select/aggregate) |
| Multiple Row Output | TableFunction (cross join/flatmap) | TableAggregateFunction (flatAggregate) |



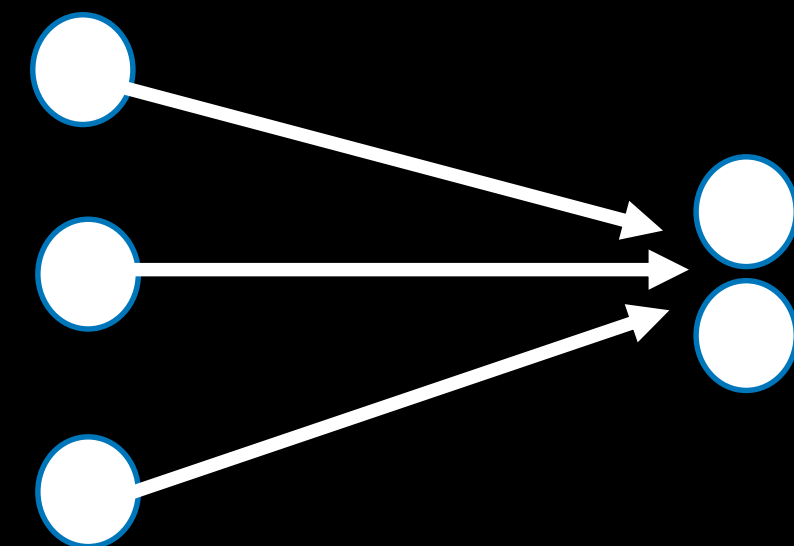
ScalarFunction
(select/map)



TableFunction
(cross join/flatmap)



AggregateFunction
(select/aggregate)



TableAggregateFunction
(flatAggregate)

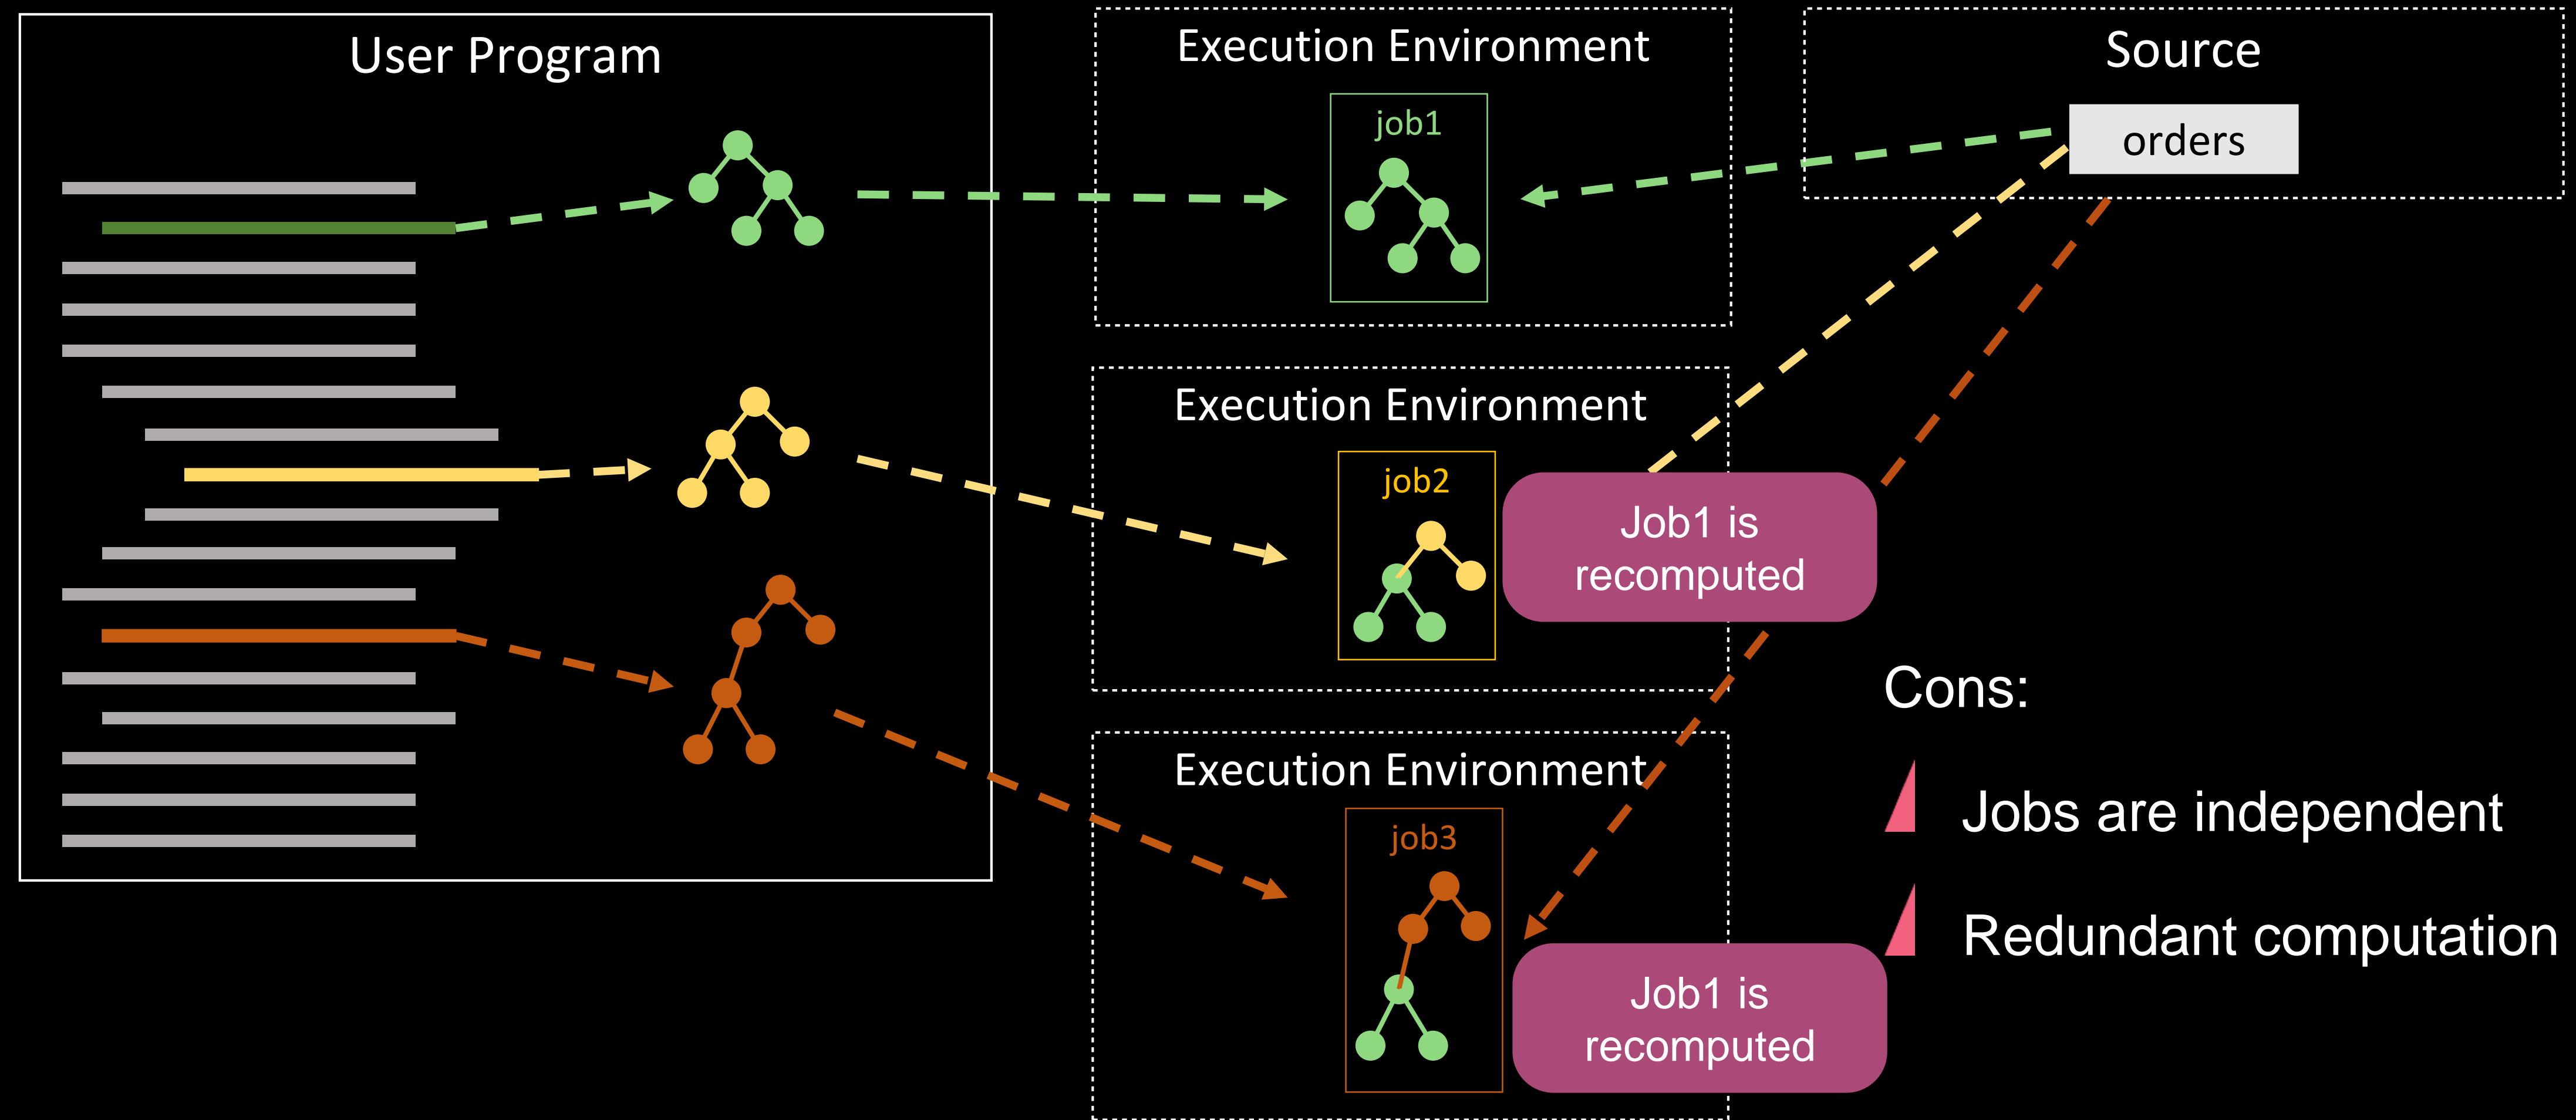
Agenda

- TableAPI enhancement: functionality and productivity
- **Interactive programming**
- Multi-language support
- Brief introduction of AI support

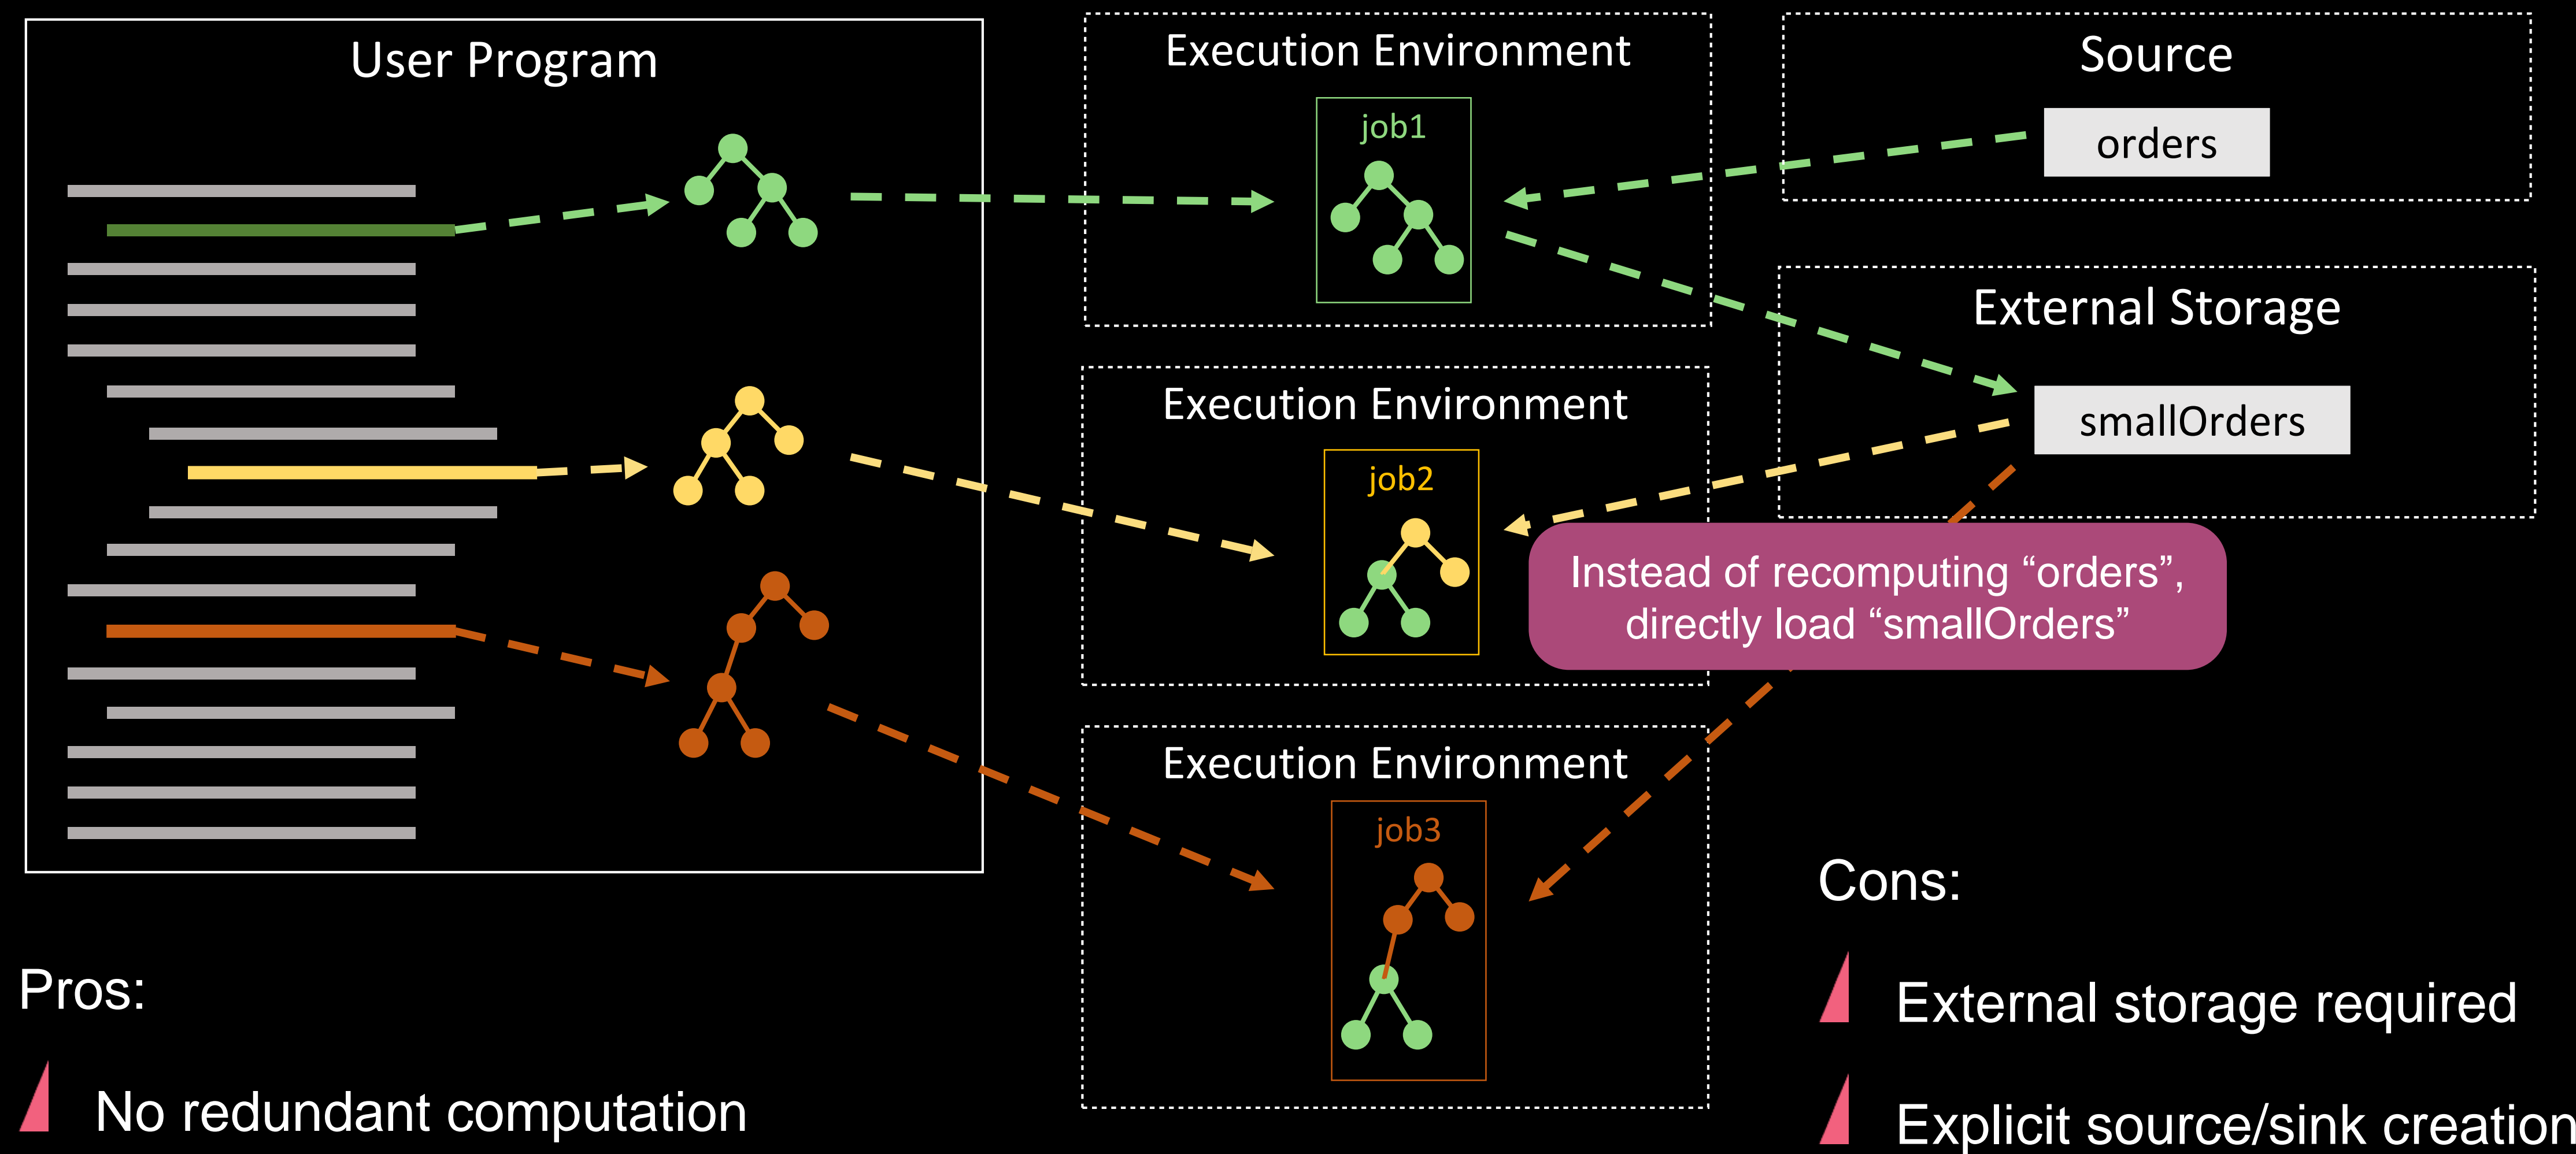
A example code snippet

```
{  
  val orders = tEnv.fromCollection(data).as ('country', 'color', 'quantity')  
  
  val smallOrders = orders.filter('quantity < 100)  
  
  val countriesOfSmallOrders = smallOrders.select('country').distinct()  
  countriesOfSmallOrders.print()  
  
  val smallOrdersByColor = smallOrders.groupBy('color').select('color', 'quantity.avg as 'avg')  
  smallOrdersByColor.print()  
}
```

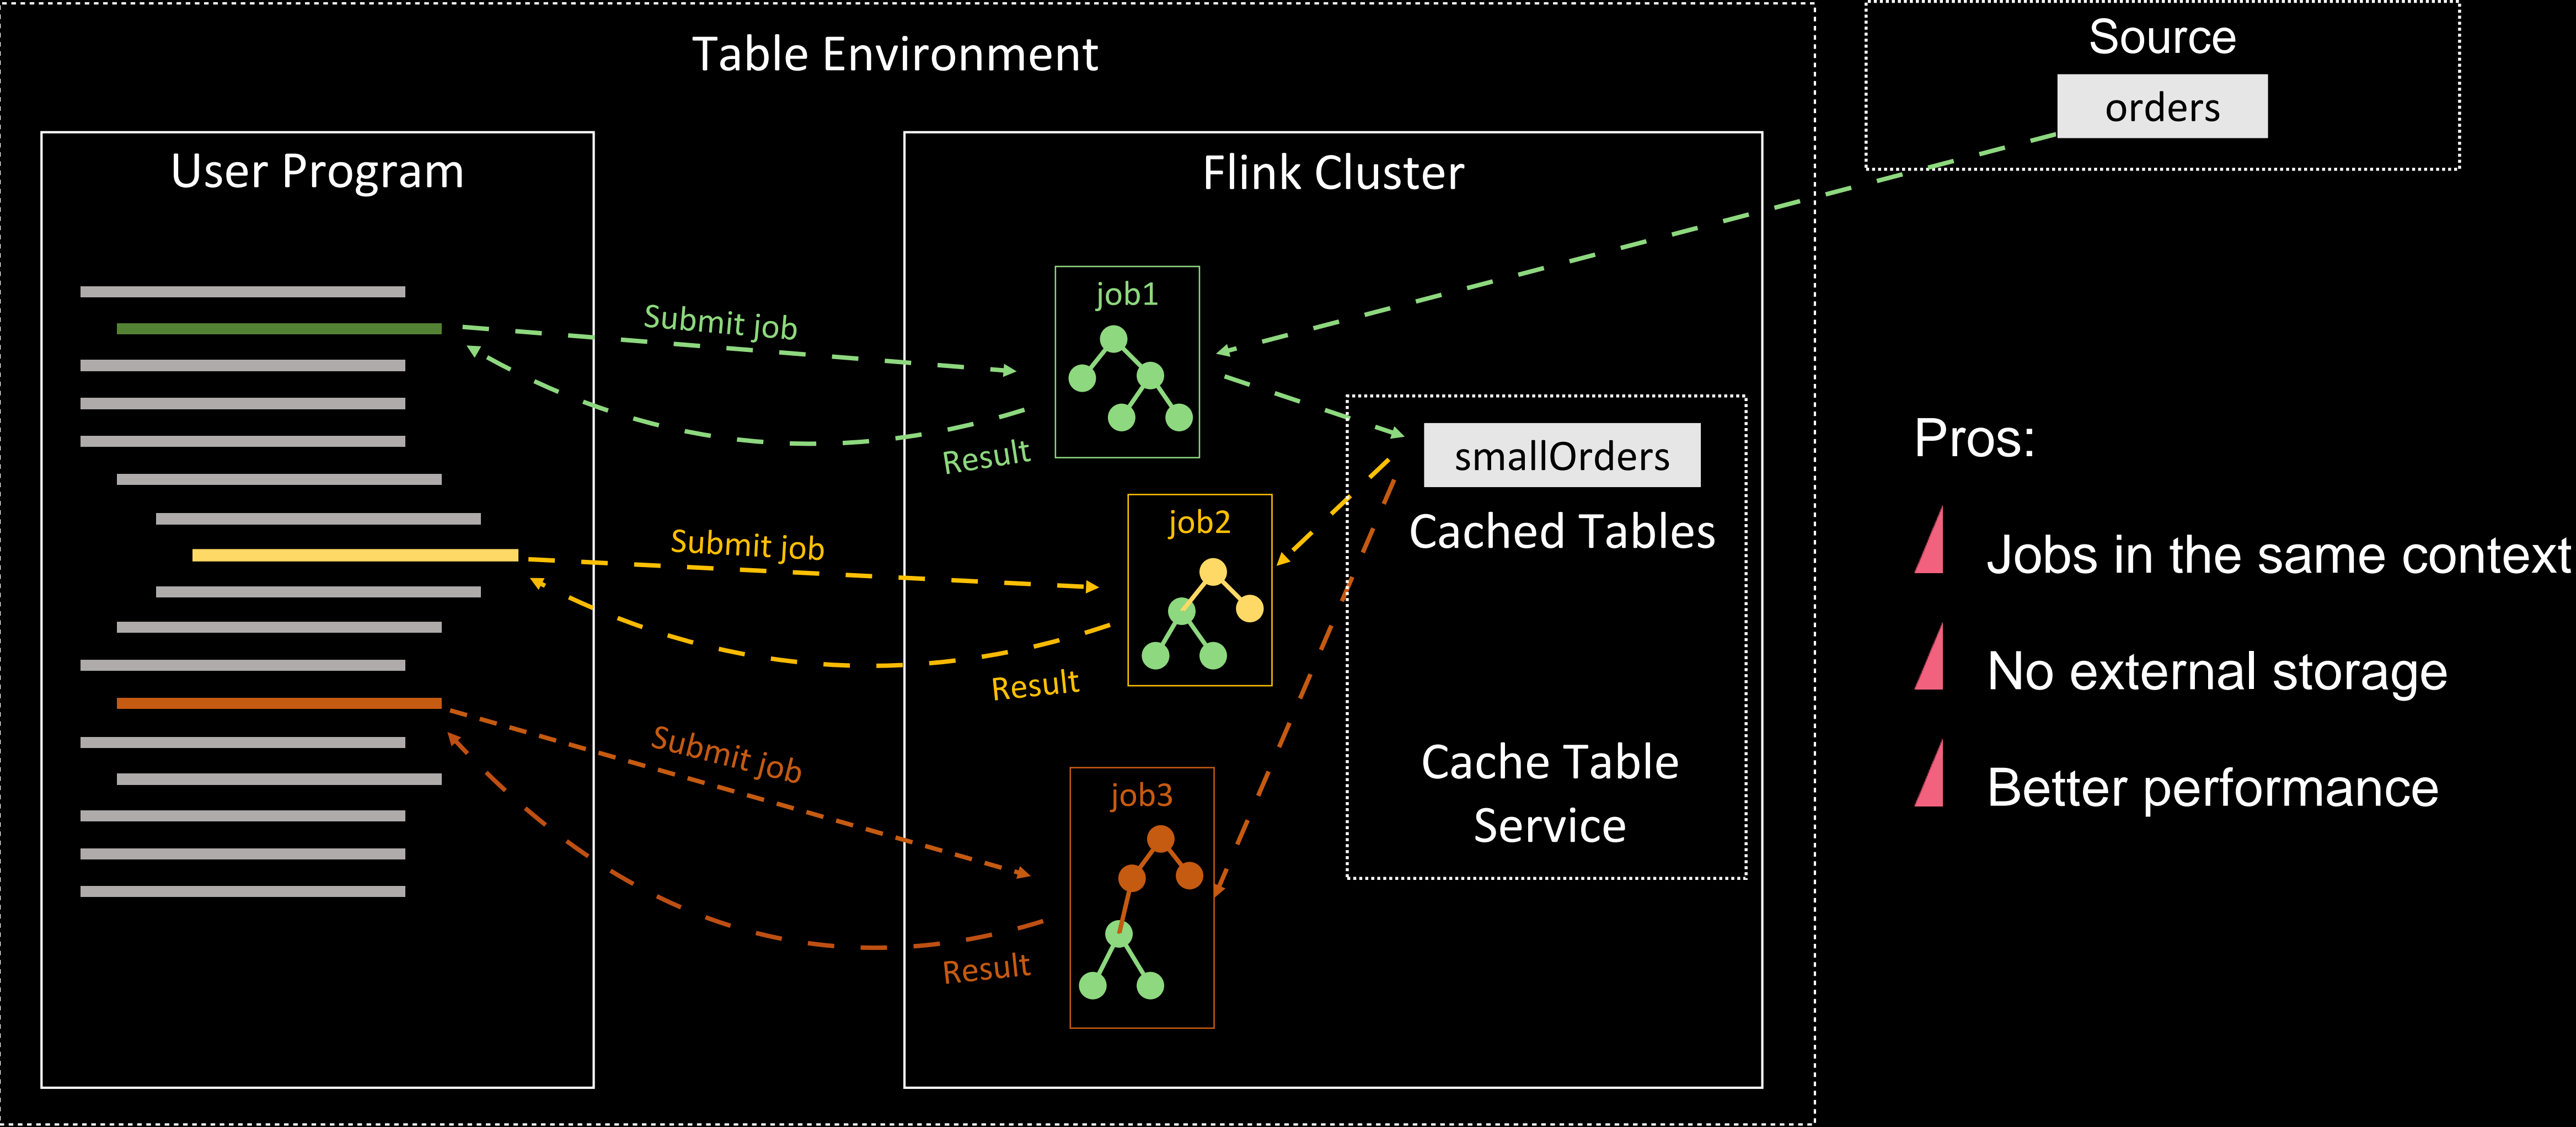
Interactive Programming without Cache or External Storage



Interactive Programming with External Storage



Interactive Programming with Cached Result (FLIP36)



Introducing Table.cache() for Interactive Programming

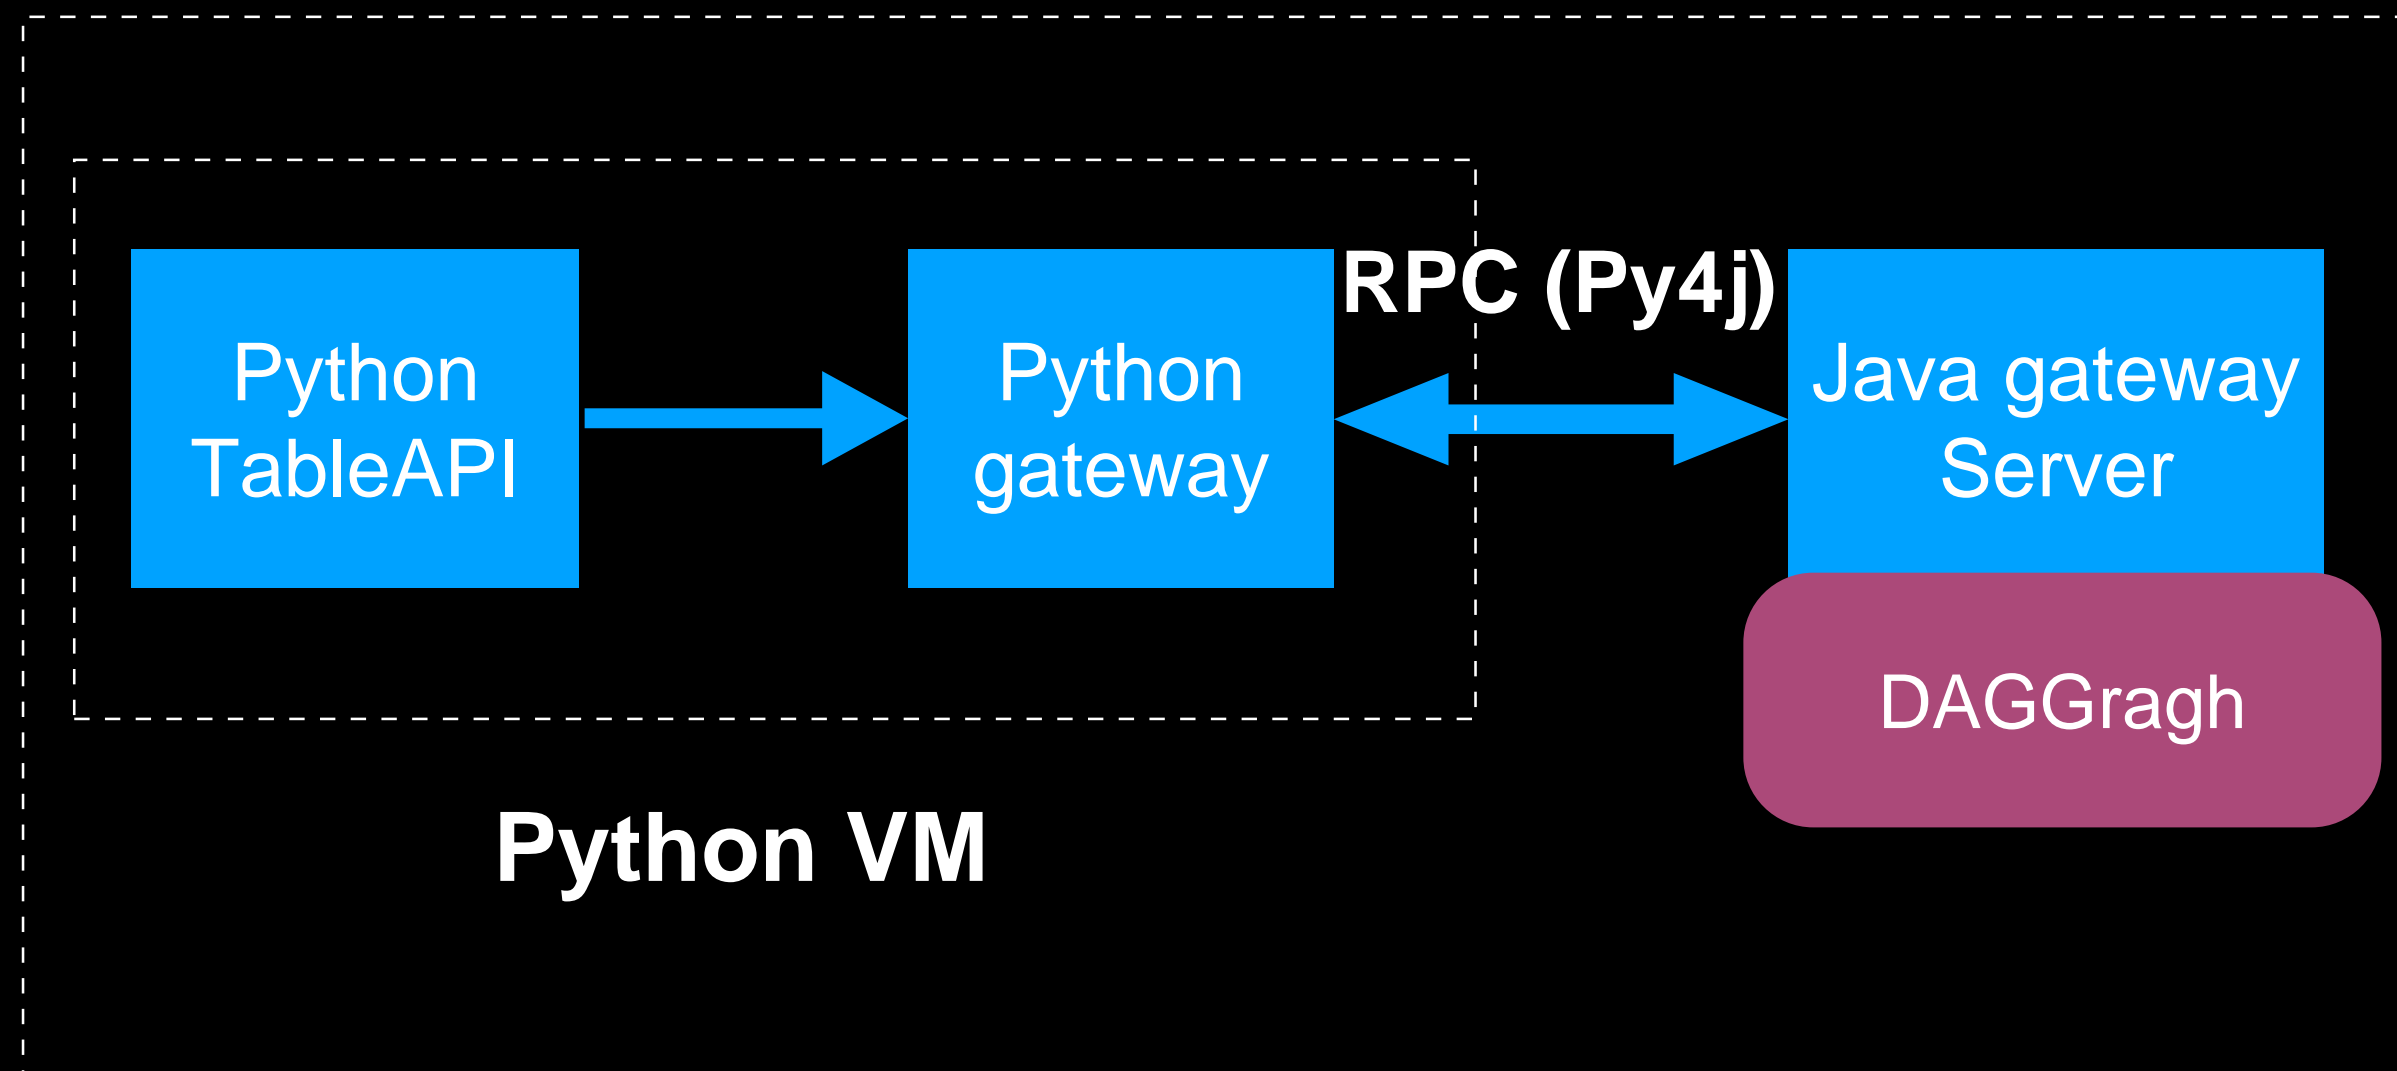
```
{  
  val orders = tEnv.fromCollection(data).as ('country', 'color', 'quantity')  
  
  val smallOrders = orders.filter('quantity < 100')  
  smallOrders.cache()  
  
  val countriesOfSmallOrders = smallOrders.select('country').distinct()  
  countriesOfSmallOrders.print()  
  
  val smallOrdersByColor = smallOrders.groupBy('color').select('color', 'quantity.avg as 'avg')  
  smallOrdersByColor.print()  
}
```

Agenda

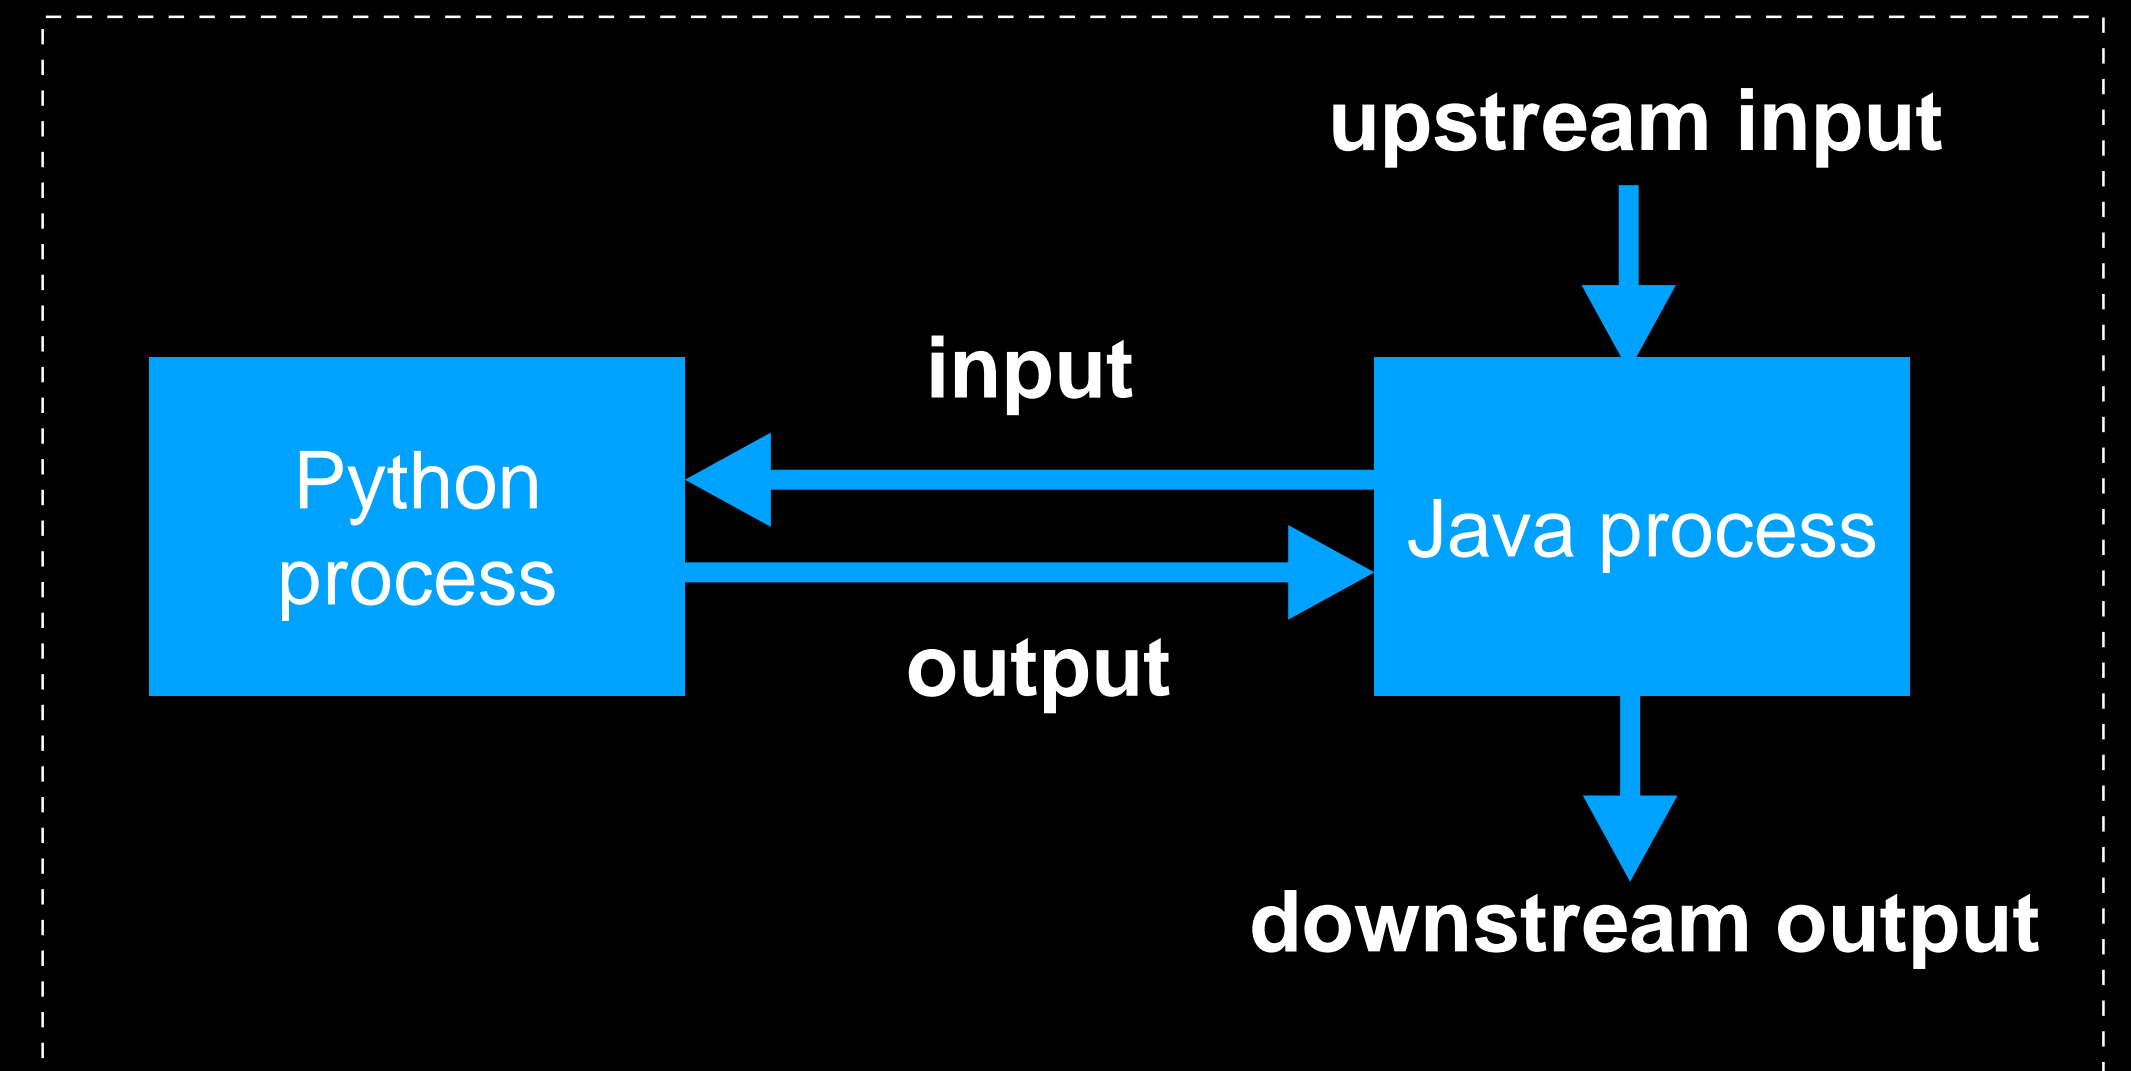
- **TableAPI enhancement: functionality and productivity**
- **Interactive programming**
- **Multi-language support**
- **Brief introduction of AI support**

Flink Python TableAPI

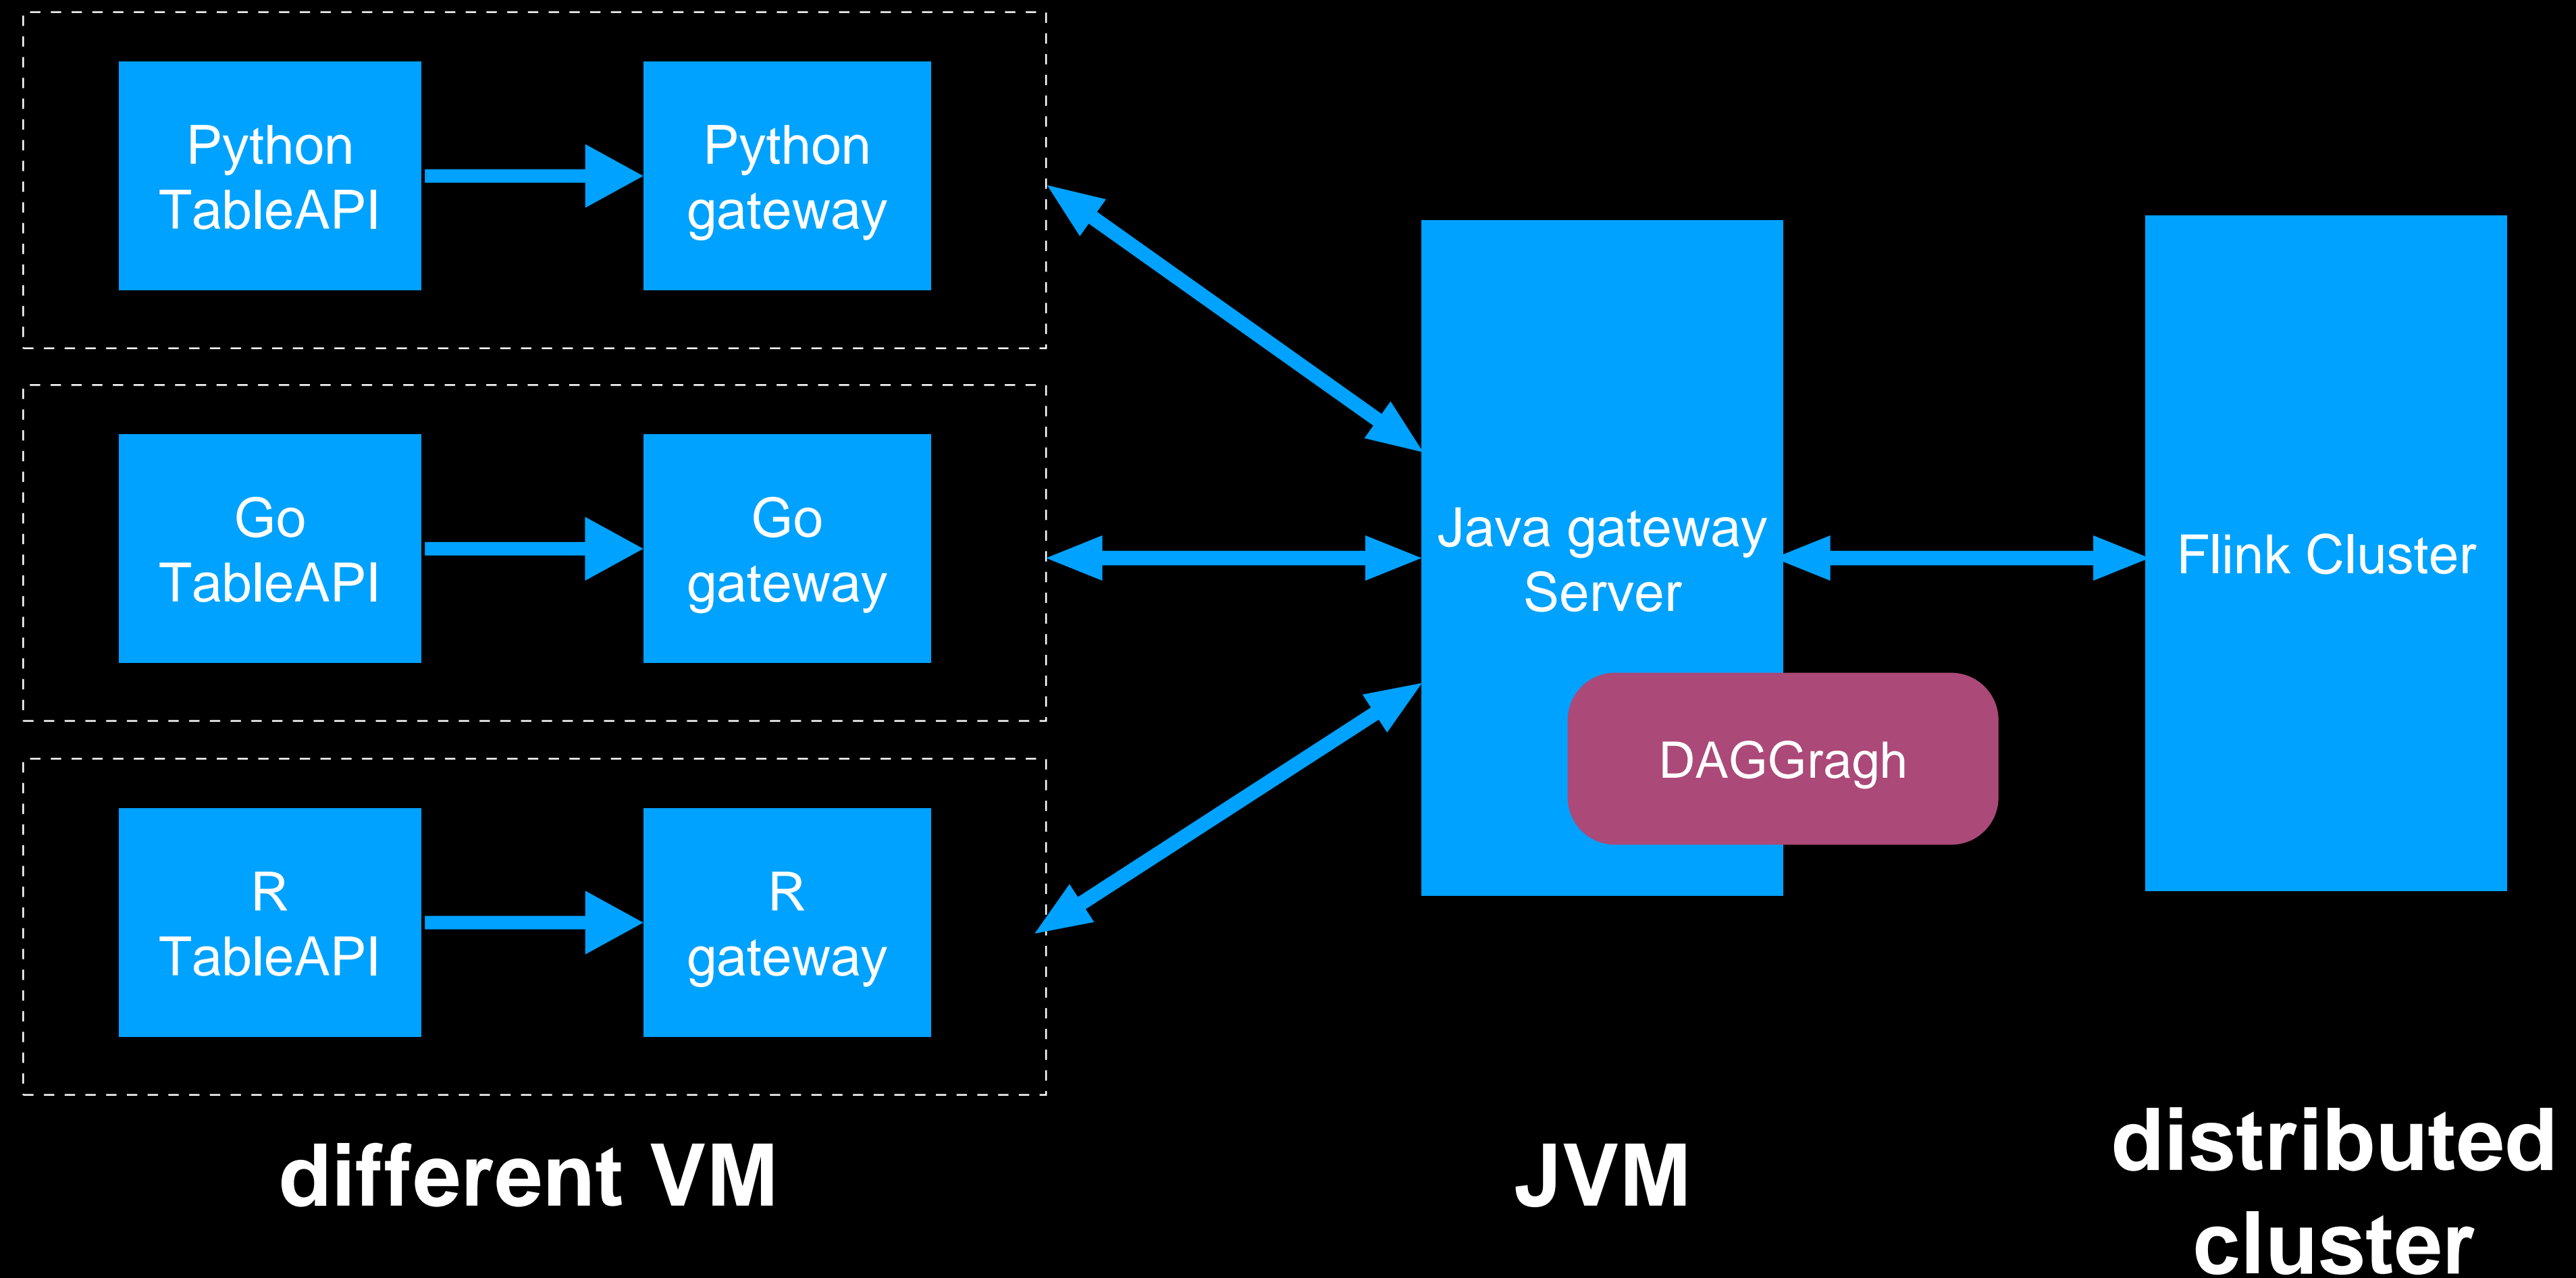
Python TableAPI



Python UDF



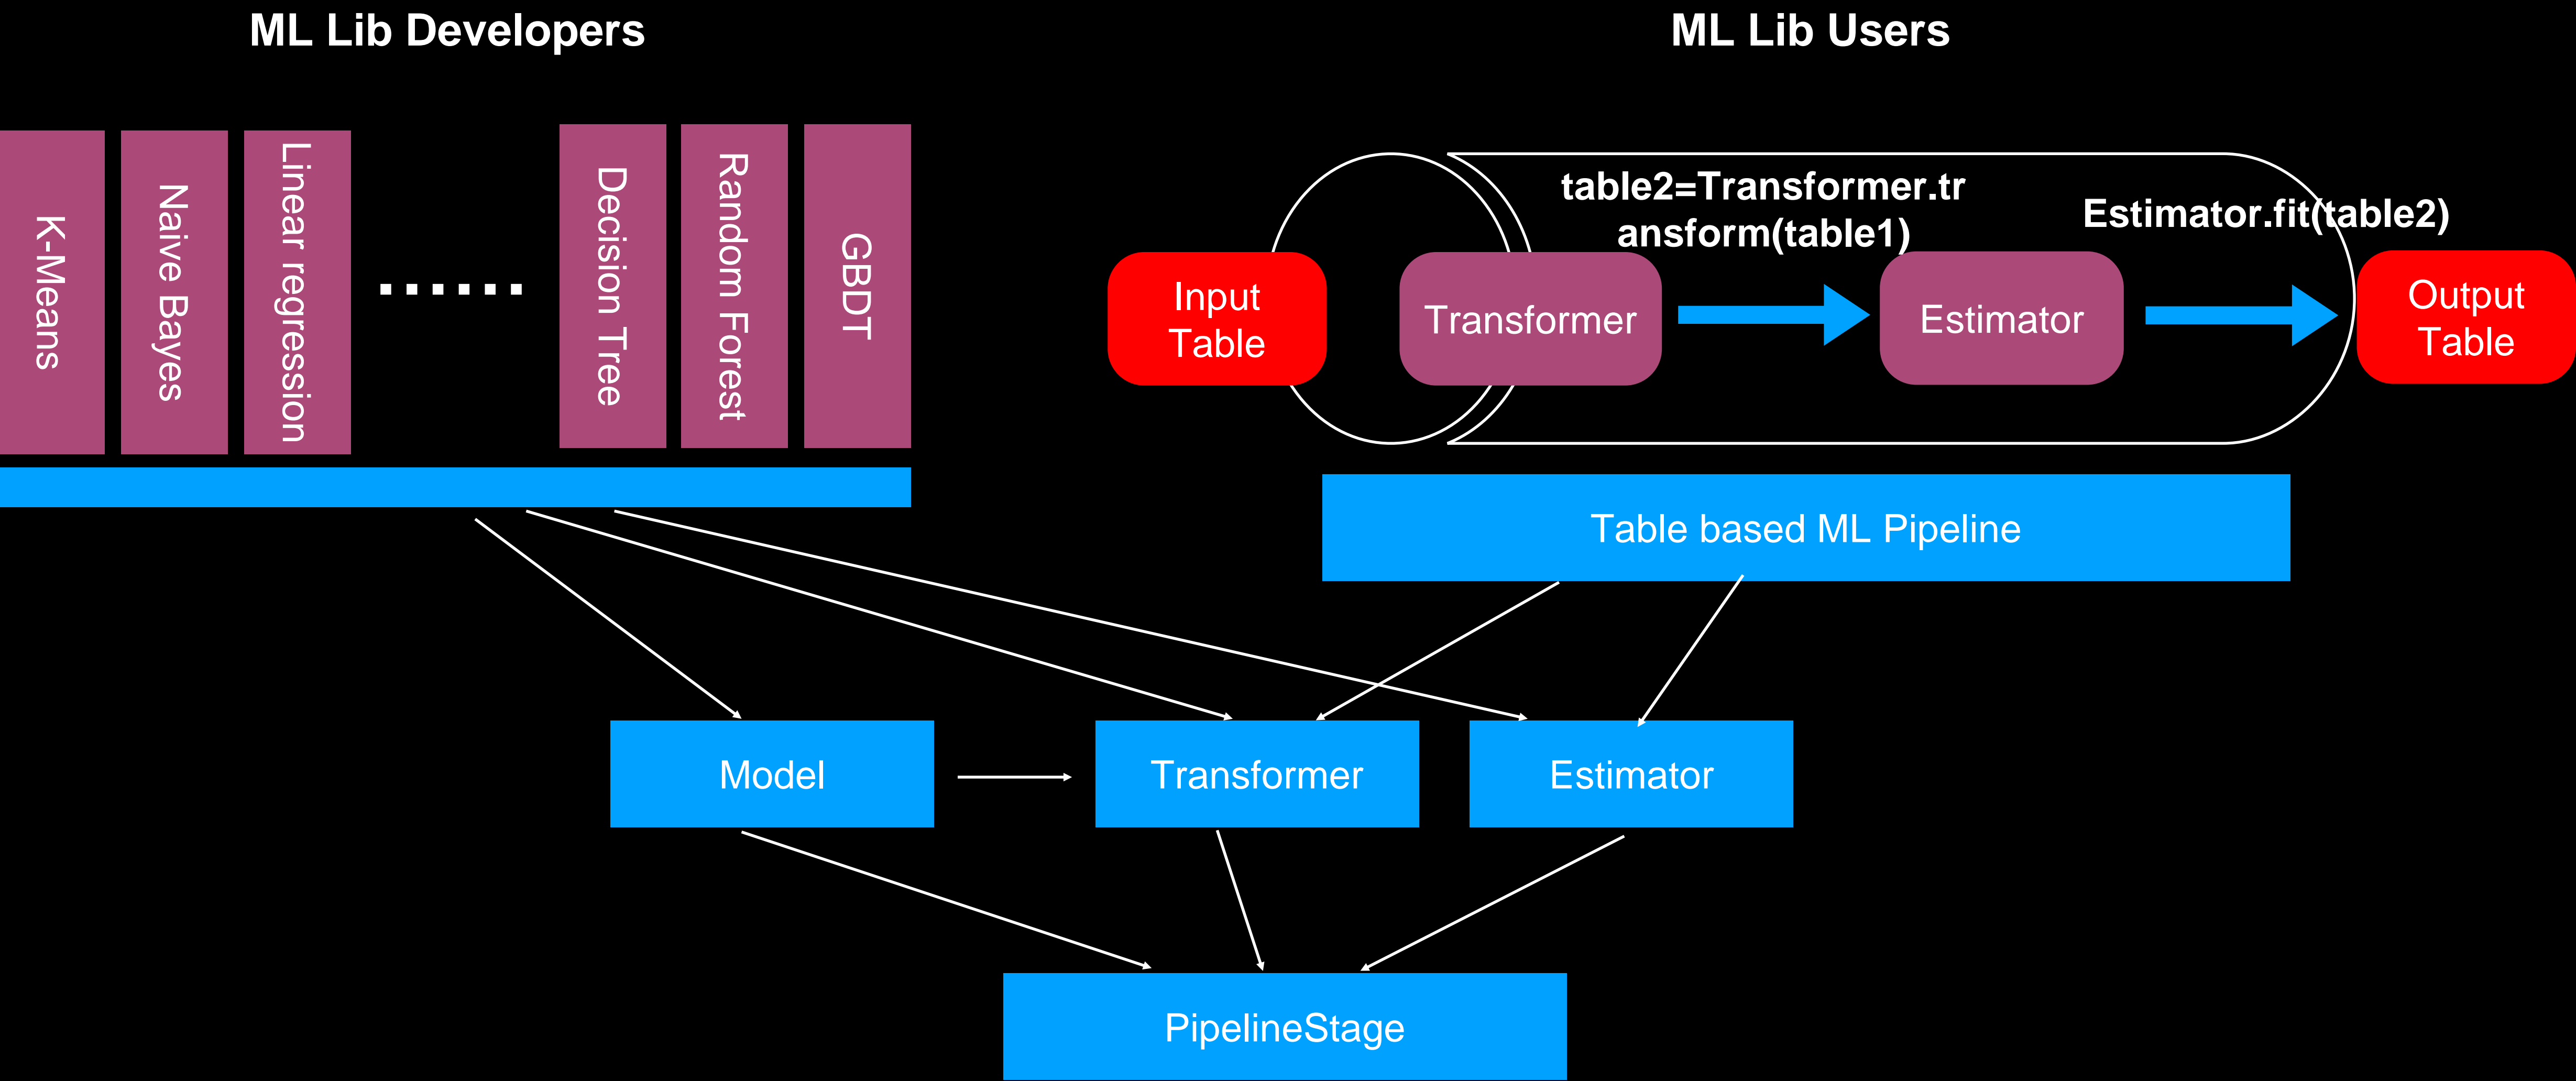
Multi-Language for Flink TableAPI



Agenda

- **TableAPI enhancement: functionality and productivity**
- **Interactive programming**
- **Multi-language support**
- **Brief introduction of AI support**

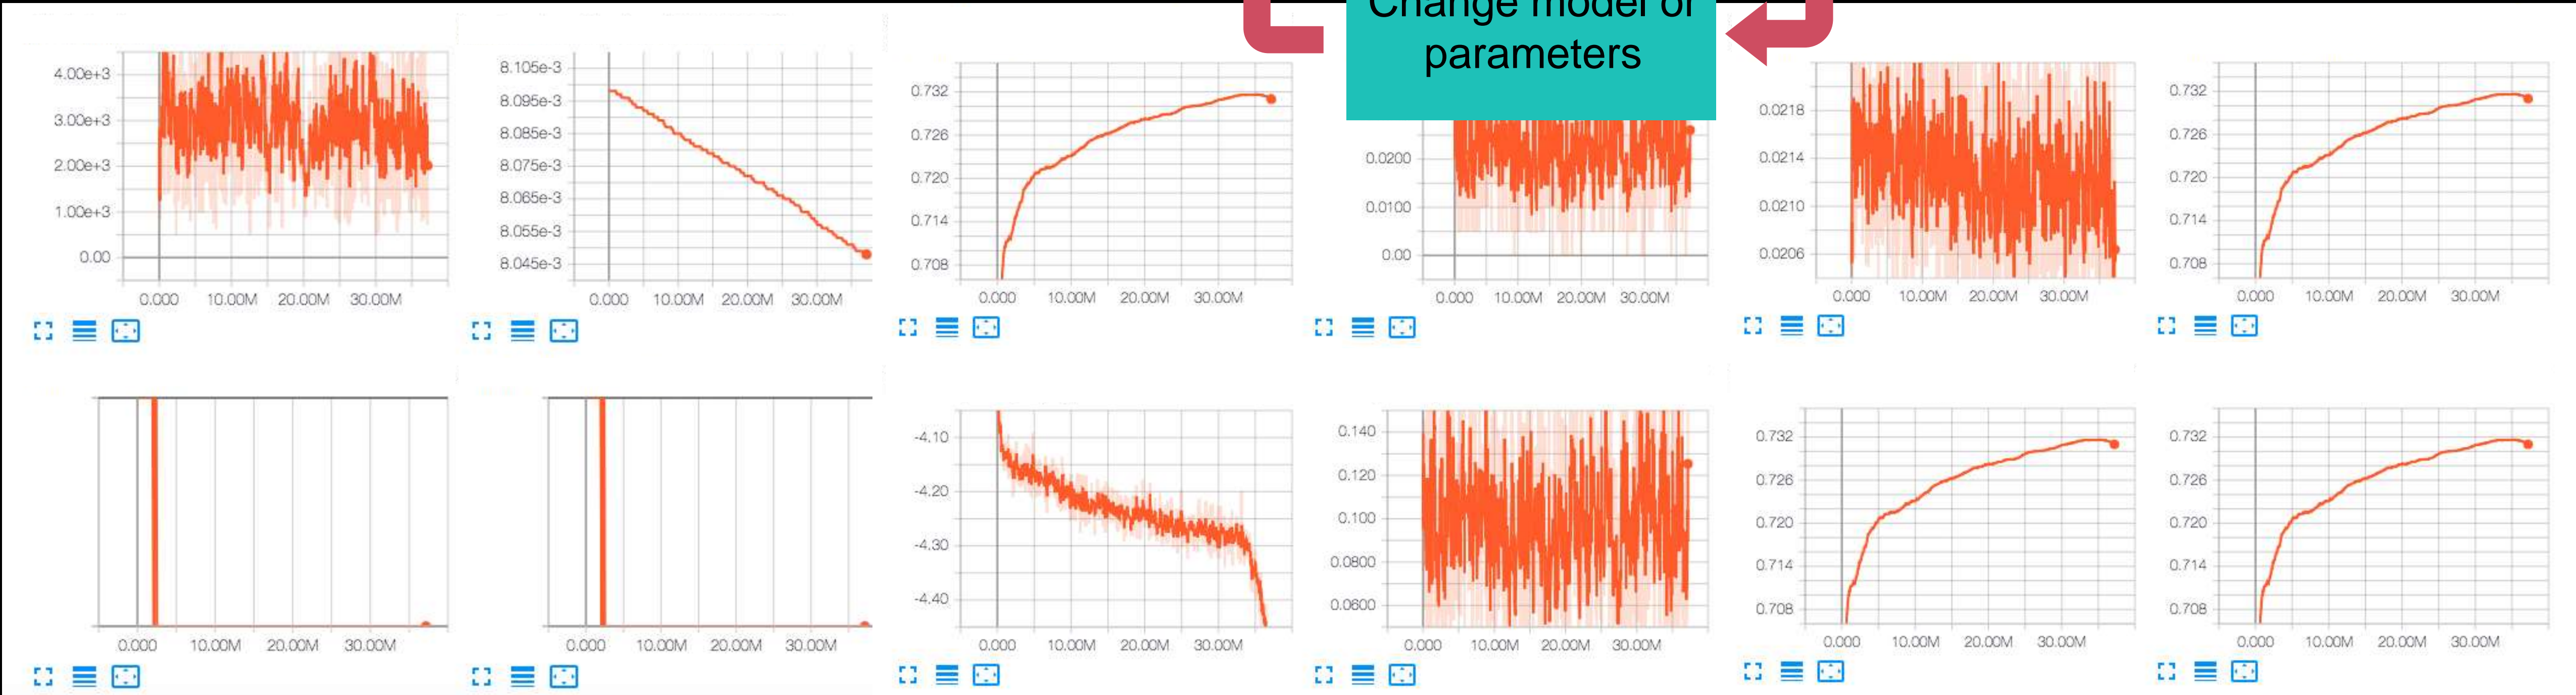
ML Pipeline - Overview (Target for Flink 1.9)



Deep Learning Pipeline on Flink



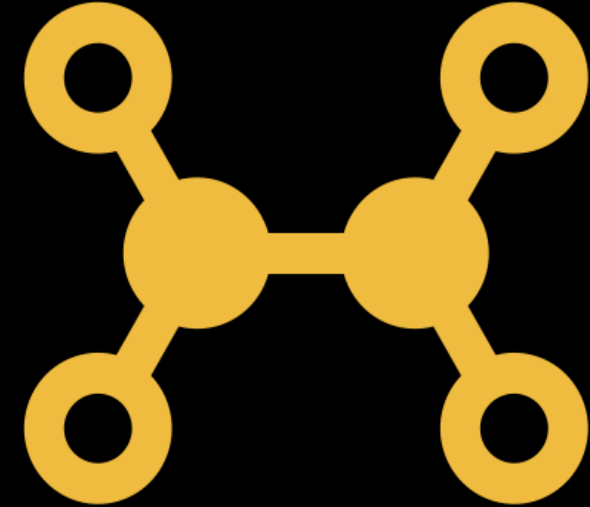
Change model or parameters



Find More Details in Another Session



Iteration



**Execute MLlib
and DL engine**

When Table meets AI: Build Flink AI Ecosystem on Table API

Shaoxuan Wang,
Director, Senior Staff Engineer at
Alibaba

4:30pm - 5:10pm

Nikko II & III

Summary

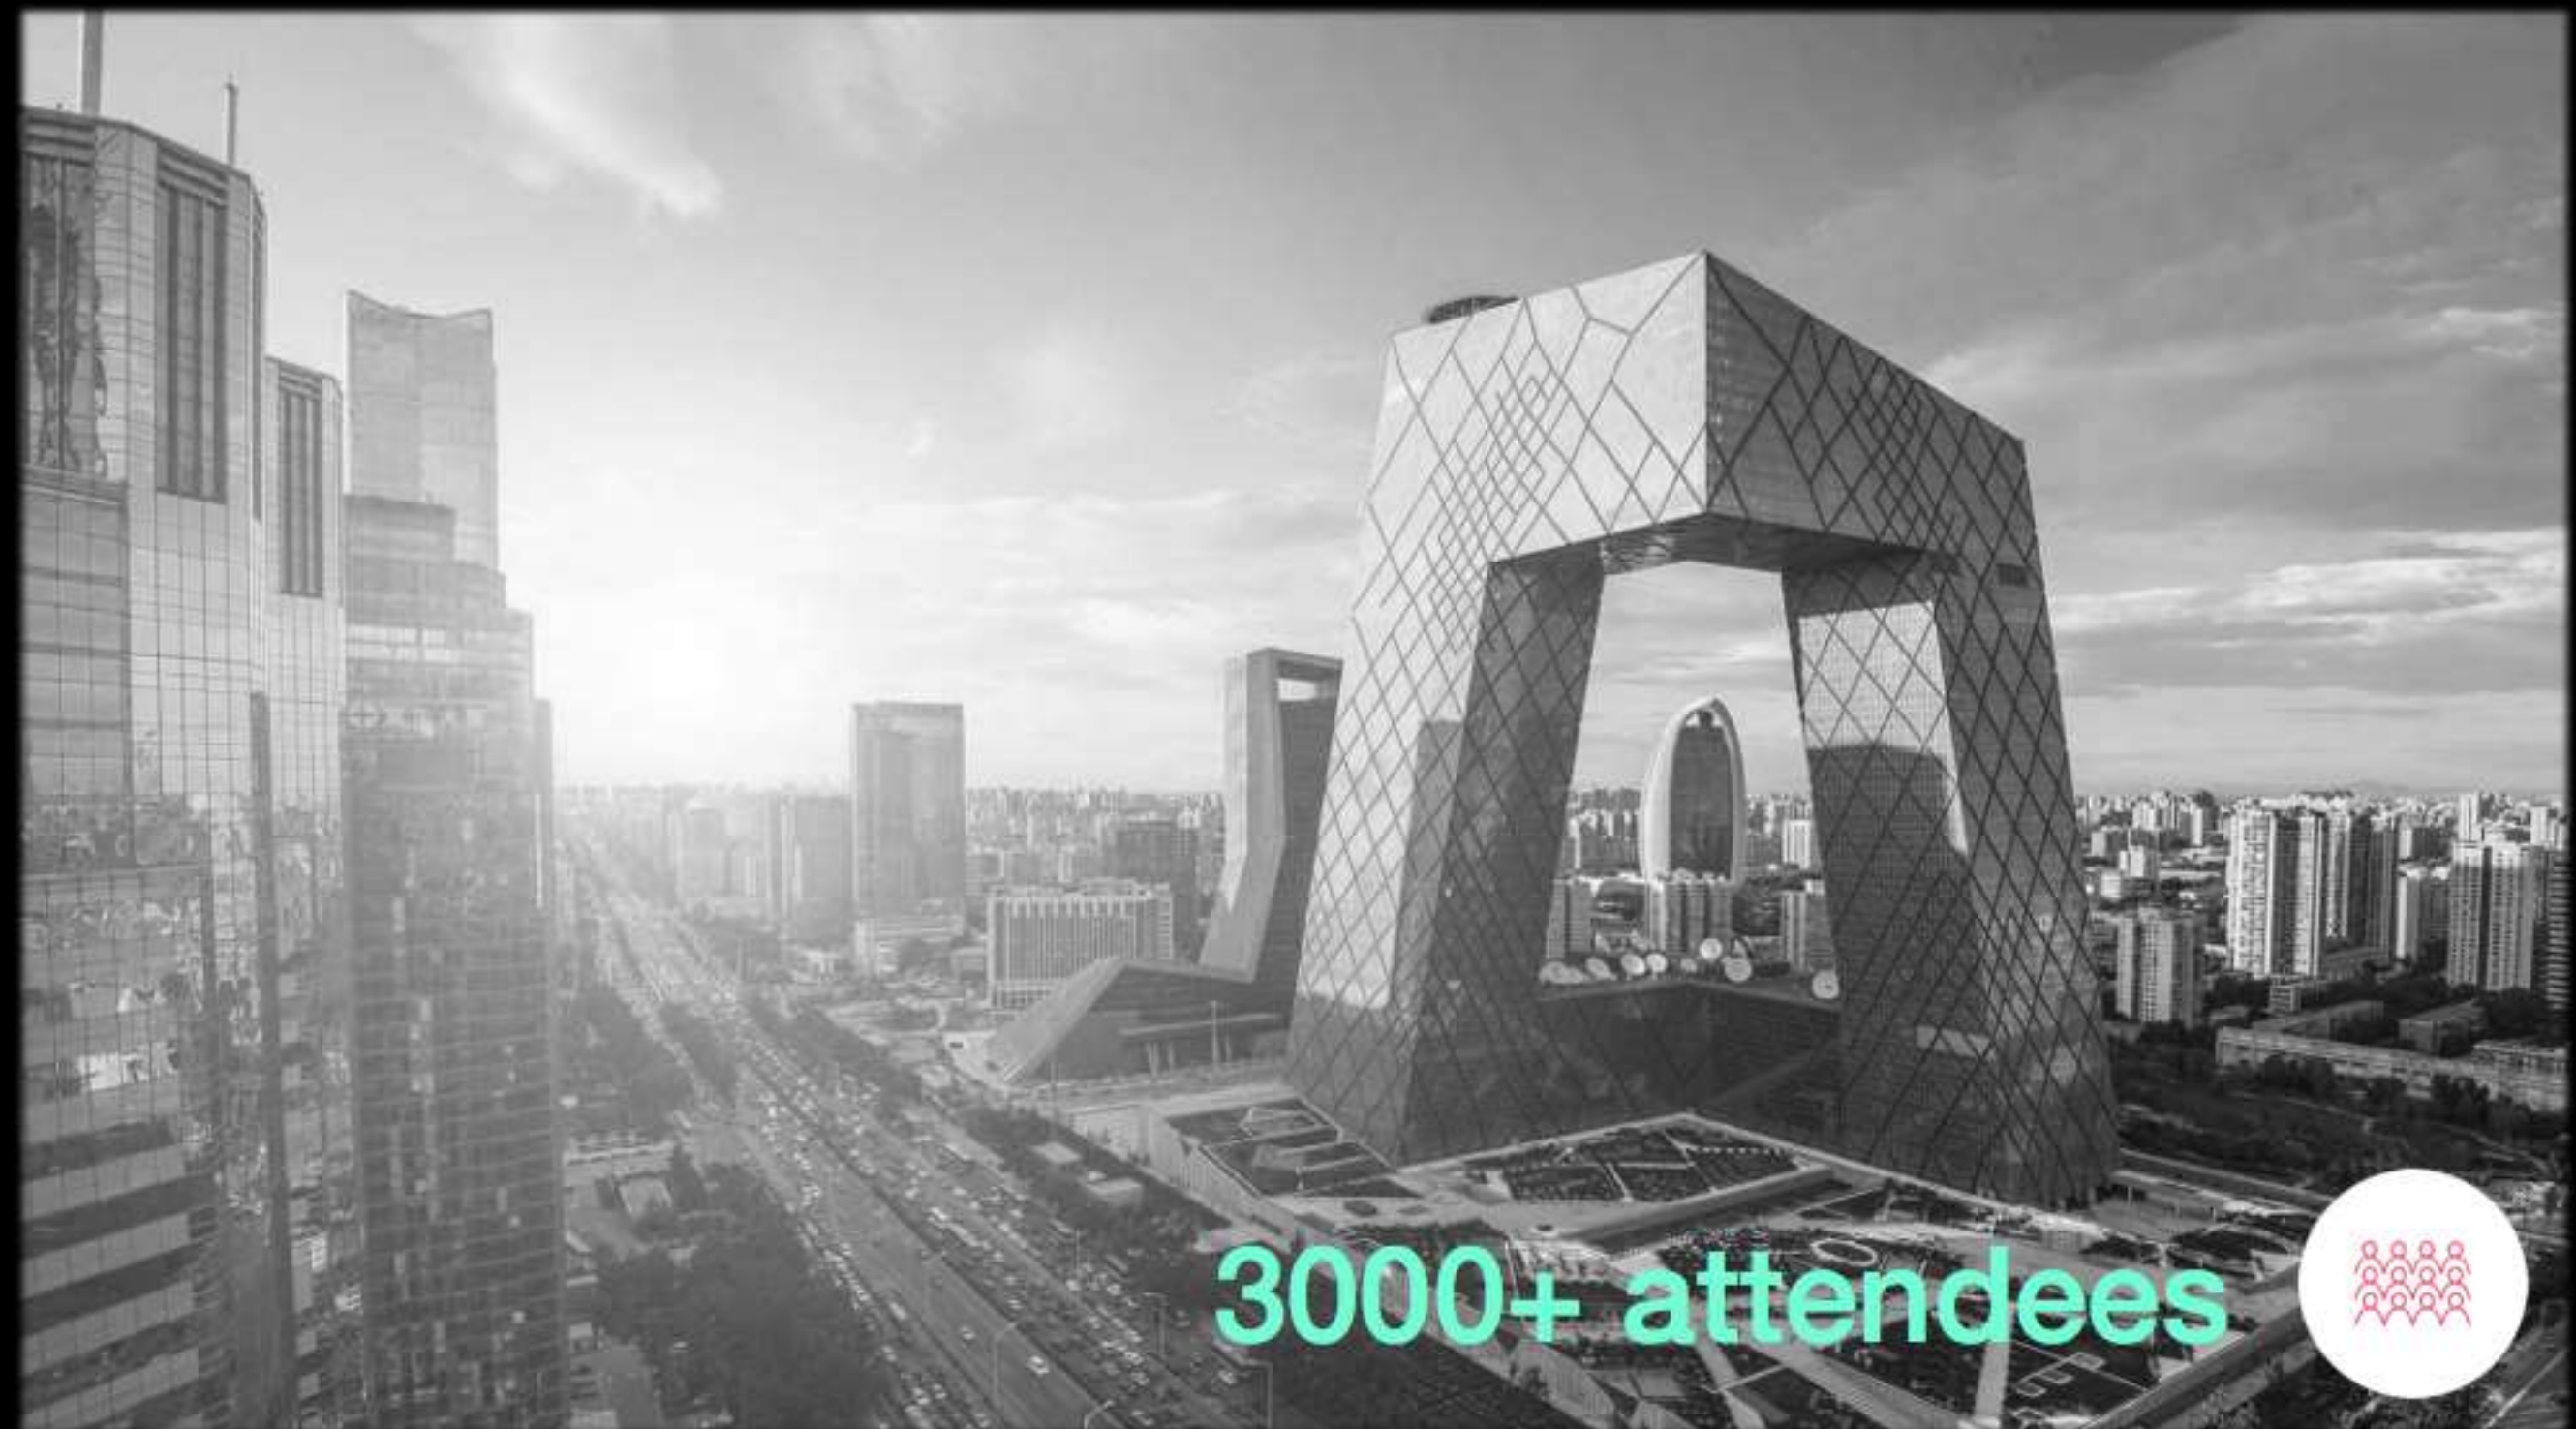
- TableAPI enhancement: functionality and productivity (FLIP29 etc.)
- Interactive programming (FLIP36, Flink1.9)
- Multi-language support (Python TableAPI, Flink1.9)
- Flink ML pipeline (Flink1.9)
- Deep learning pipeline on Flink (Will be open-sourced)

Welcome to Flink Forward - ASIA

Dec. 2019 @ Beijing

flink-forward-china@list.alibaba-inc.com

Q & A



3000+ attendees

