



Tencent 腾讯

Developing and Operating Real-Time Applications at Tencent

Xiaogang SHI
robbieshi@tencent.com



PART I
Background

PART II
**Introduction to
Oceanus**

PART III
**Improvement to
Flink**

About Tencent

2

- The largest social communities in China in terms of MAU
- Weixin & Wechat: 1 billion MAU
- QQ: 807 million MAU

**SOCIAL
NETWORK**

- The leading mobile payment platform by active users and transactions in China.
- 900 million MAU
- 1 billion transactions daily

FINTECH

- The leading online video streaming platform in China in terms of mobile DAU and subscriptions
- 89 million subscriptions
- Up 58% year-over-year

VIDEO

- The largest online music platform in China
- 800 million MAU

MUSIC

- The top mobile news app by active users
- 150 million MAU

NEWS

MEDIA

- Covered 25 regions and operated 53 availability zones

CLOUD


Tencent
腾讯

GAMES


- The largest gaming company in the world by active users and revenue
- 200 million MAU

Real-Time Applications at Tencent

3



ETL



Monitoring



Real-Time BI



Online Learning

210 Million

Maximum number of
messages received
per second

17 Trillion

Number of messages
received per day

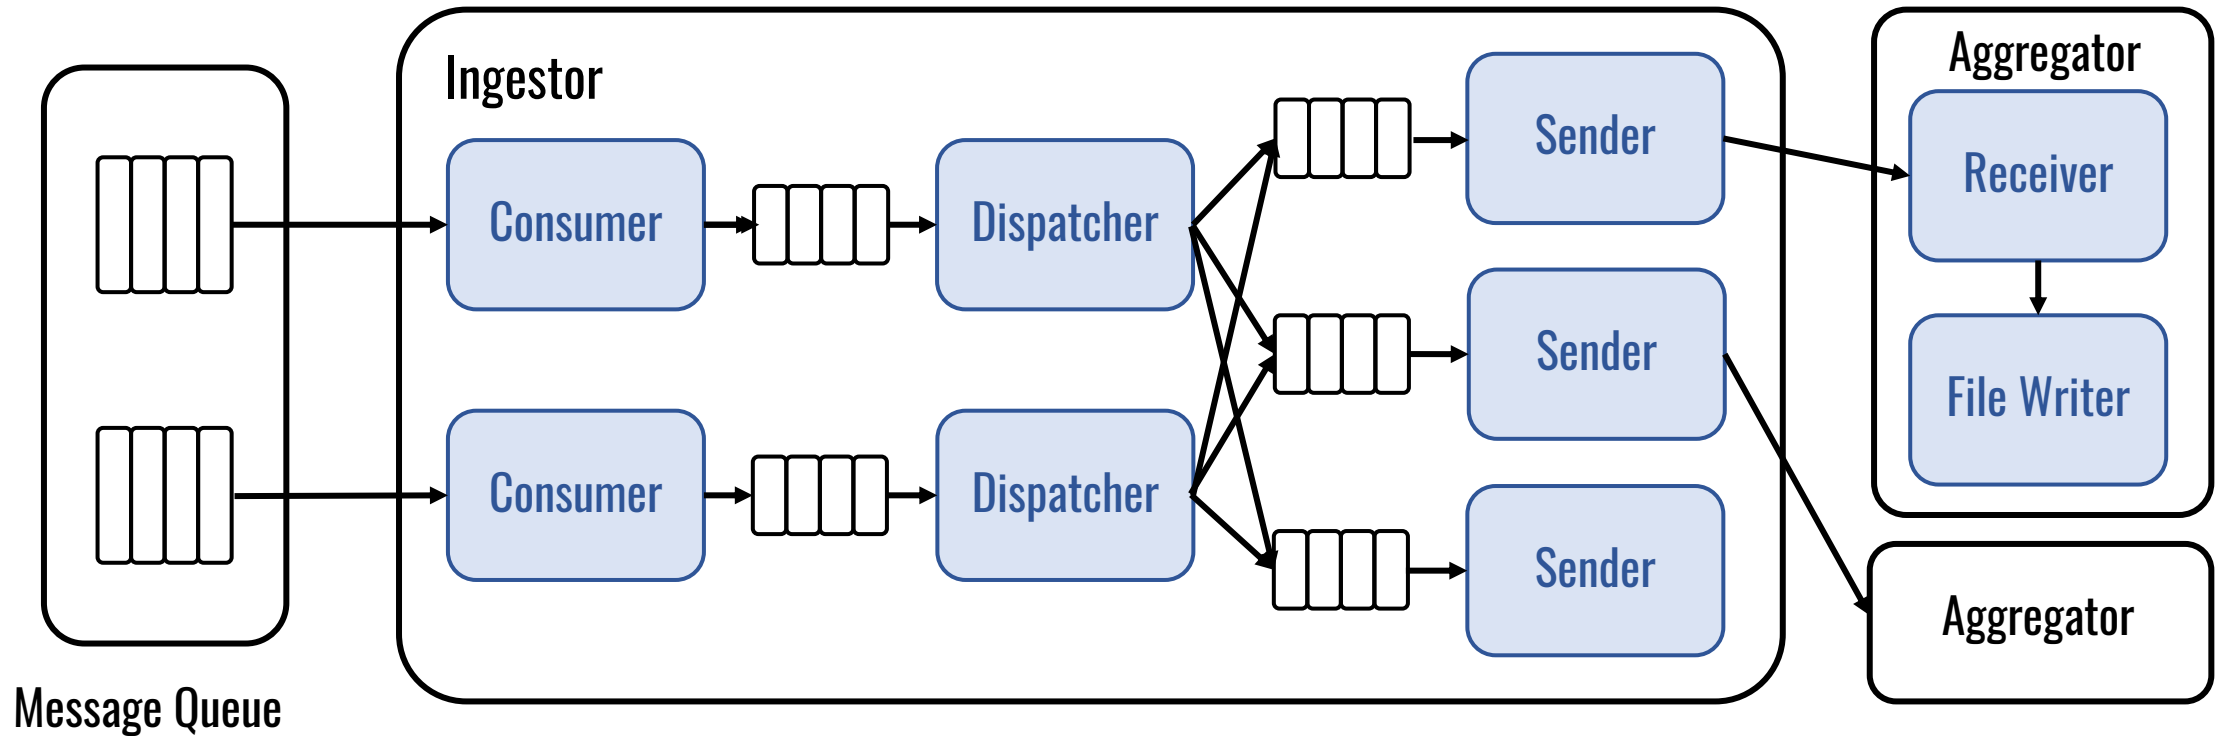
3 PB

Amount of data
received per day

Why Flink?

4

An ETL pipeline as an example



It's very difficult to improve performance while ensuring correctness.

Flink's Strength

5



Efficient support for states
Automatic distribution while
rescaling



AT-LEAST and EXACTLY-
ONCE guarantees while
tolerating failures



Flexible and powerful
programming interfaces for
stream processing

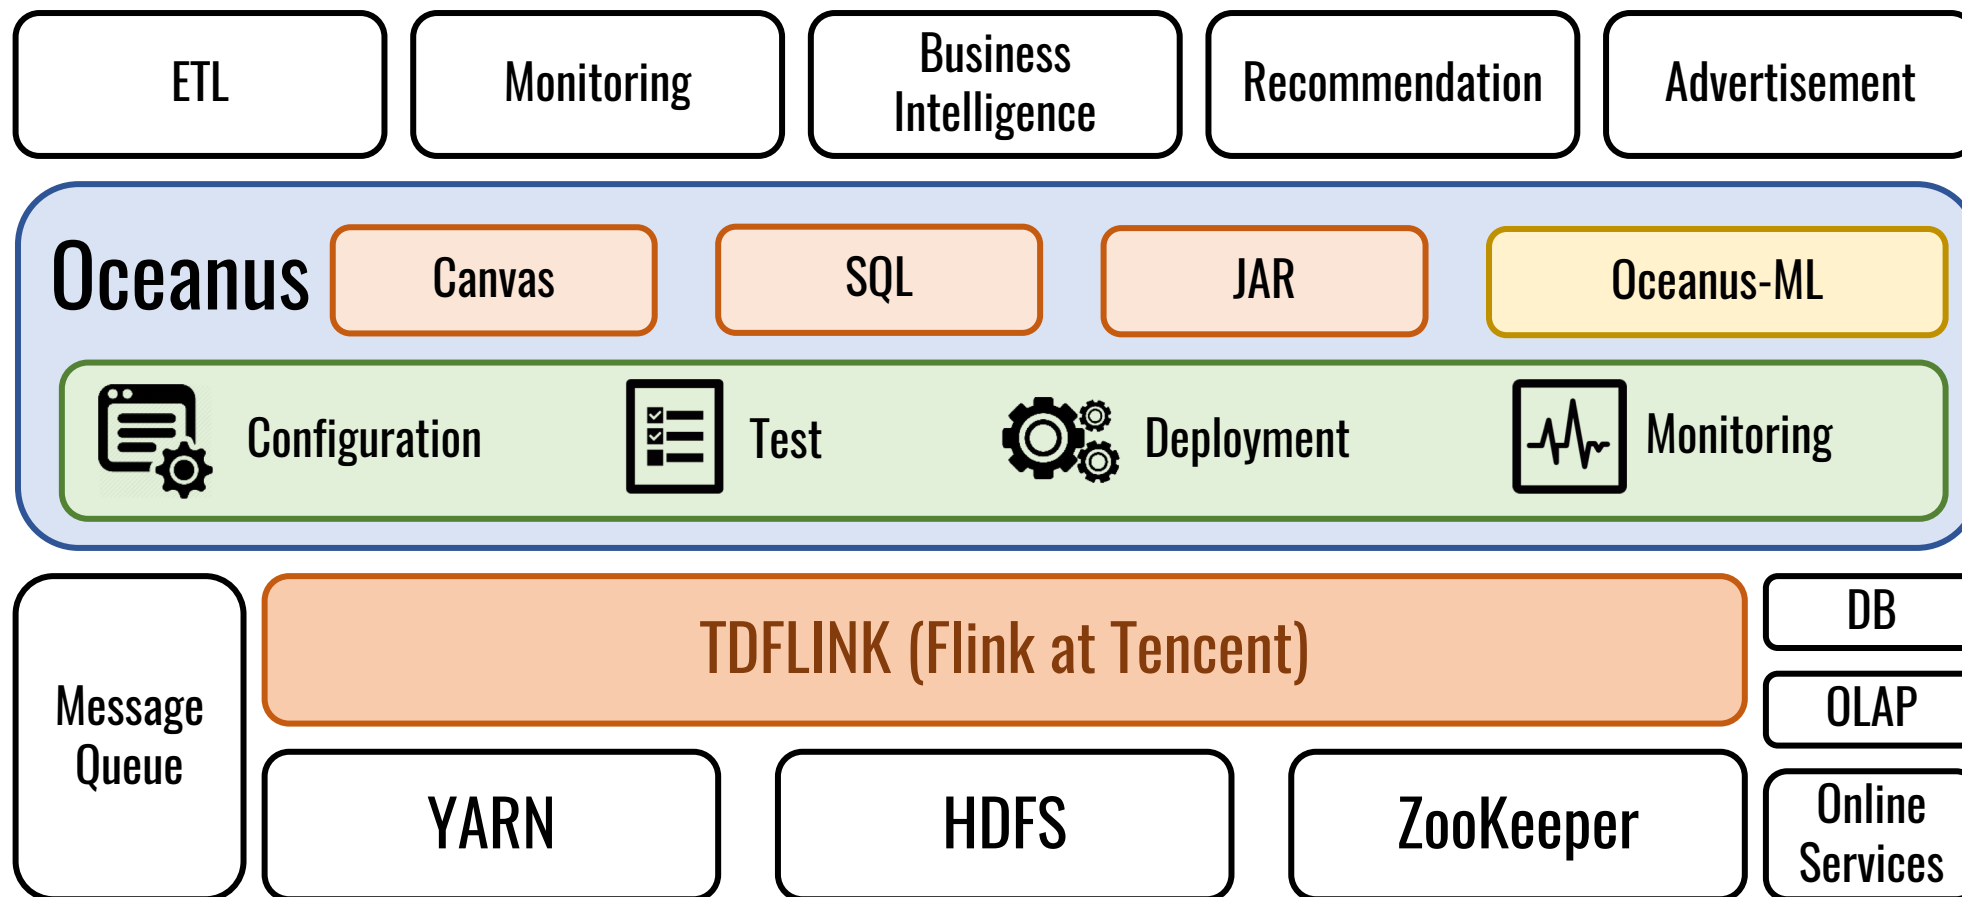


Low latency and high
throughput

Oceanus Overview


6

A unified platform to develop and operate real-time applications



Developing with Oceanus

7



Type:	Kafka	
Name:	test_topic	
Fields:	Name	String
	Id	Integer
	Etime	Long
	Phone	String

Type:	Tumbling	
Length:	10	seconds

Users can easily develop their applications by dragging and connecting operators.


Developing with Oceanus

8




Compile

Submit



Jobs are managed with refined Cluster Clients.

```
INSERT INTO uni_hot_stat (  
  song_id,  
  batch_time,  
  qy_play_score  
)  
SELECT  
  song_id,  
  sum(play_score) AS qy_play_score  
FROM kfk_u002  
GROUP BY song_id
```



JobGraph

Users can easily configure their applications with visualized job graphs.

Operating with Oceanus

Improve operating efficiency with rich metrics.

Job Details									
Vertices		Metrics (Records)		Metrics (Bytes)					
ID	Name	Watermark max ▼	InQueue Usage max ▼	OutQueue Usage max ▼	Input Records sum ▼	Output Records sum ▼	Input Rate sum ▼	Output Rate sum ▼	
0	Source: Collection Source bc764cd8ddf7a0cff126f51c16239658	N/A	0.00%	100.00%	0	3,425,953,507,895	0	3,527,340	
1	Map -> Sink: Unnamed 20ba6b65f97481d5570070de90e4e791	N/A	0.00%	0.00%	3,425,953,469,803	0	3,527,350	0	

1. An operator is bottleneck if its in-queue is full while its out-queue is not full.
2. The ratio between the throughput of different operators remain roughly the same when the parallelism changes. We can utilize this property to configure the parallelism (Don't work well with window operators).
3. There may exist data skew when the difference between the maximum and the minimum throughput is very large.

Much information can be obtained from thread stacks.

Checkpoint timeout: lock unreleased by blocked user functions, slow hdfs writes, blocked user checkpoint functions

Performance issues

Metrics	Threads	Log	Stdout	
ID	Name	CPU ↓	State	Stack
70	Map -> Sink: Unnamed (1/1)	99.77%	RUNNABLE	org.apache.flink.streaming.api.operators.AbstractStreamOperator.setKeyContextElement(AbstractStreamOperator.java:625) org.apache.flink.streaming.api.operators.AbstractStreamOperator.setKeyContextElement1(AbstractStreamOperator.java:61) org.apache.flink.streaming.runtime.io.StreamInputProcessor.processInput(StreamInputProcessor.java:201) org.apache.flink.streaming.runtime.tasks.OneInputStreamTask.run(OneInputStreamTask.java:107) org.apache.flink.streaming.runtime.tasks.StreamTask.invoke(StreamTask.java:301) org.apache.flink.runtime.taskmanager.Task.run(Task.java:721) java.lang.Thread.run(Thread.java:748)
68	Source: Collection Source (1/1)	80.32%	RUNNABLE	org.apache.flink.runtime.state.KeyGroupRangeAssignment.assignToKeyGroup(KeyGroupRangeAssignment.java:59) org.apache.flink.runtime.state.KeyGroupRangeAssignment.assignKeyToParallelOperator(KeyGroupRangeAssignment.java:48) org.apache.flink.streaming.runtime.partition.KeyGroupStreamPartitioner.selectChannels(KeyGroupStreamPartitioner.java:101) org.apache.flink.streaming.runtime.partition.KeyGroupStreamPartitioner.selectChannels(KeyGroupStreamPartitioner.java:101) org.apache.flink.runtime.io.network.api.writer.RecordWriter.emit(RecordWriter.java:101) org.apache.flink.streaming.runtime.io.StreamRecordWriter.emit(StreamRecordWriter.java:81) org.apache.flink.streaming.runtime.io.RecordWriterOutput.pushToRecordWriter(RecordWriterOutput.java:107) org.apache.flink.streaming.runtime.io.RecordWriterOutput.collect(RecordWriterOutput.java:89) org.apache.flink.streaming.runtime.io.RecordWriterOutput.collect(RecordWriterOutput.java:45) org.apache.flink.streaming.api.operators.AbstractStreamOperator\$CountingOutput.collect(AbstractStreamOperator.java:7) org.apache.flink.streaming.api.operators.AbstractStreamOperator\$CountingOutput.collect(AbstractStreamOperator.java:7) org.apache.flink.streaming.api.operators.StreamSourceContexts\$NonTimestampContext.collect(StreamSourceContexts.java:43) org.apache.flink.streaming.api.functions.source.FromIteratorFunction.run(FromIteratorFunction.java:43) org.apache.flink.streaming.api.operators.StreamSource.run(StreamSource.java:94) org.apache.flink.streaming.api.operators.StreamSource.run(StreamSource.java:58) org.apache.flink.streaming.runtime.tasks.SourceStreamTask.run(SourceStreamTask.java:99) org.apache.flink.streaming.runtime.tasks.StreamTask.invoke(StreamTask.java:301) org.apache.flink.runtime.taskmanager.Task.run(Task.java:721) java.lang.Thread.run(Thread.java:748)
62	SystemResourcesCounter probing thread	0.19%	TIMED_WAITING	java.lang.Thread.sleep(Native Method) org.apache.flink.runtime.metrics.util.SystemResourcesCounter.run(SystemResourcesCounter.java:128)

Improvement to Flink

11

Job Management

- ✓ Avoid split-brain with Zookeeper transactions
- ✓ Non-disruptive recovery of job masters
- Fine-grained recovery of tasks with cached result partitions

Resource Management

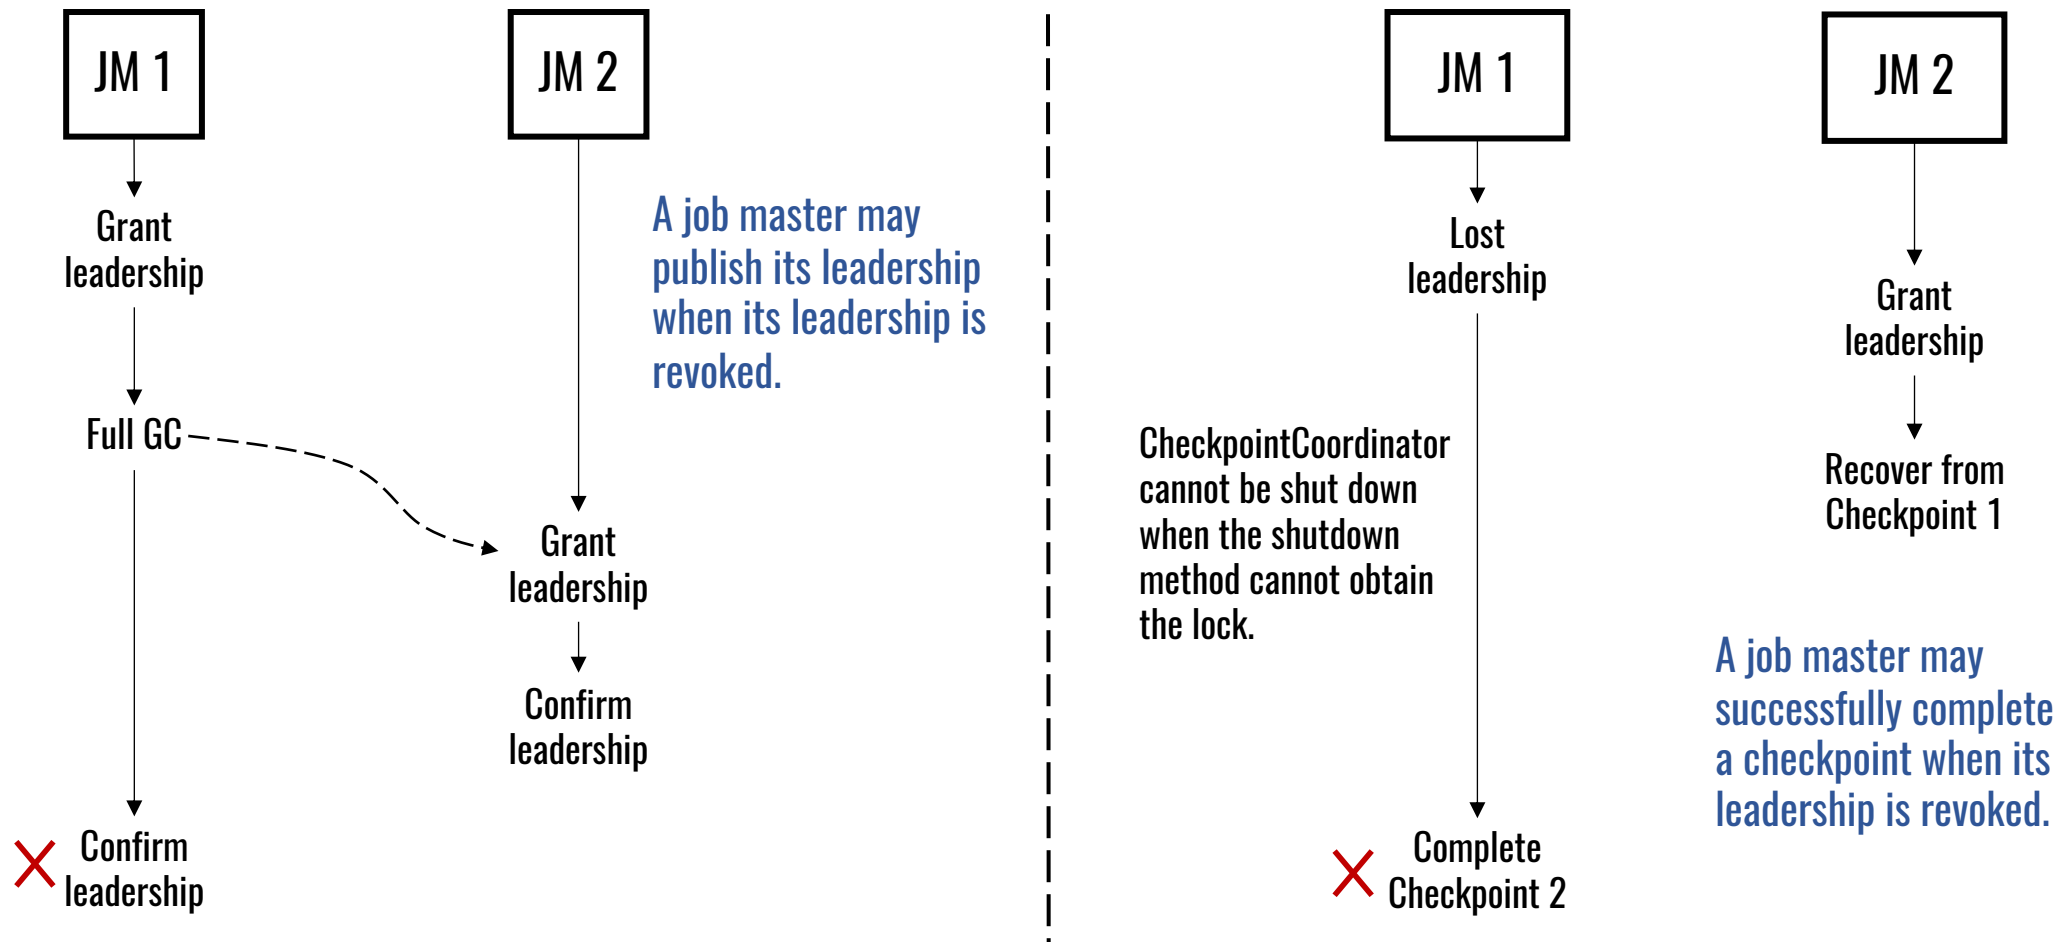
- ✓ Fine-grained resource allocation
- Improve scheduling efficiency

Performance & Usability

- ✓ Local keyed streams
- ✓ Incremental windows
- ✓ UDX
- ✓ DimJoin
- ✓ Top N

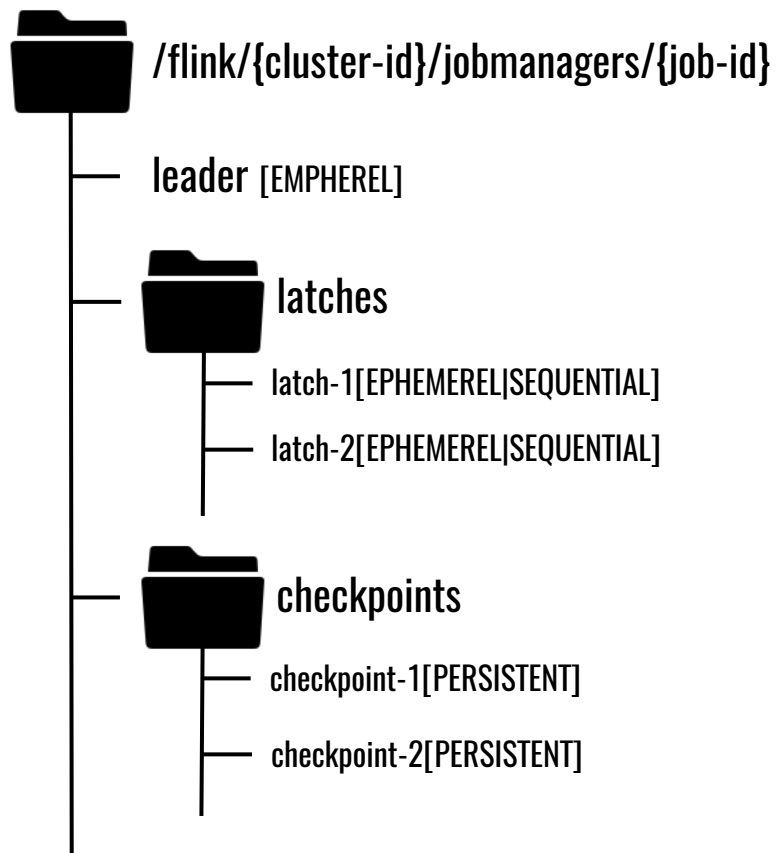
Refine leader coordination

Current problems: It's difficult to reason about leadership in Flink.



Refine leader coordination

13



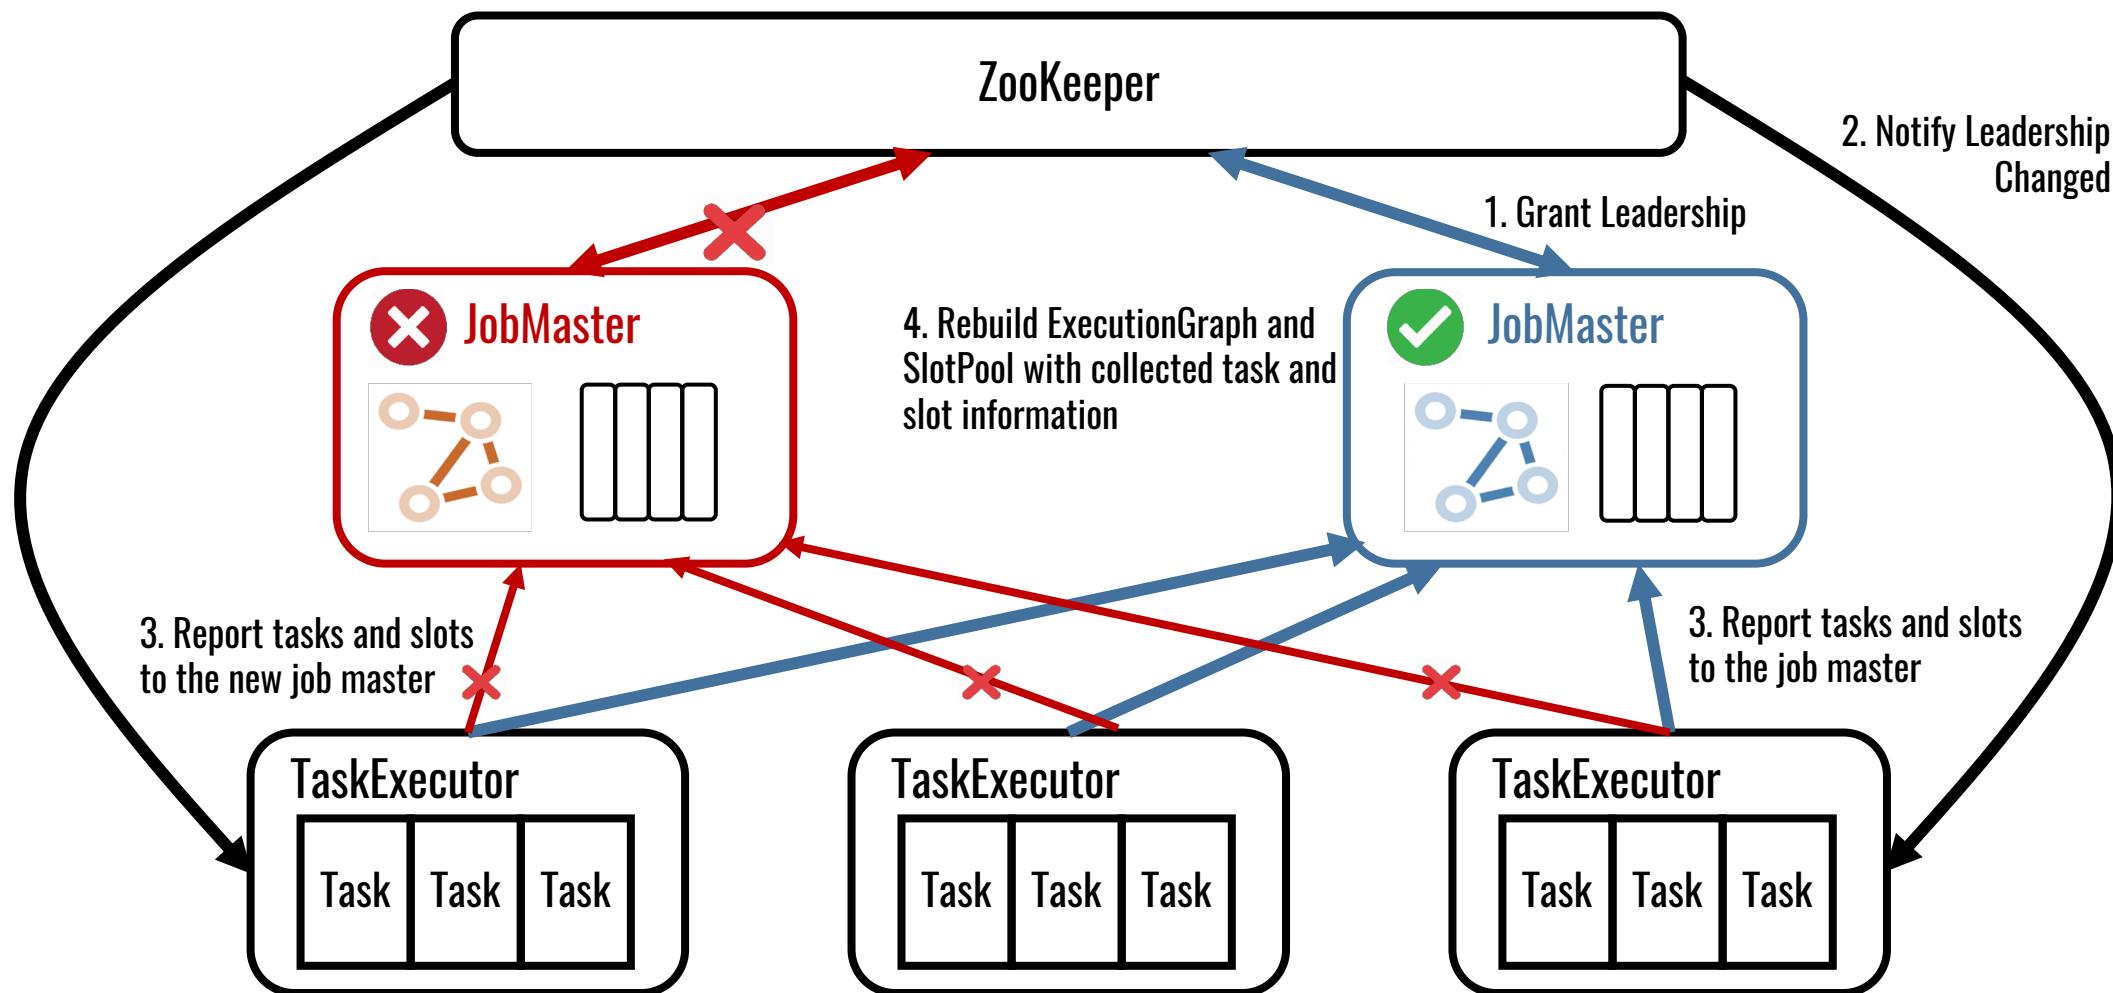
- Each leader contender creates a EPHEMERAL and SEQUENTIAL latch.
- The contender whose latch's sequential number is smallest is elected as the leader.
- A leader's leadership is granted as long as its latch exists.
- Each contender can only access states when it has granted leadership and its latch still exists.

```
zkClient.inTransaction()  
    .check().forPath(myLatch).and()  
    .setData().forPath(dataPath).and()  
    .commit()
```

Non-disruptive recovery of job masters

14

Avoid restarting tasks when the job master fails.




Fine-grained resource allocation

15

Current Problems:

- The resource specification for operators does not take effect in resource allocation.
- Slots are allocated according to the number of available slots in task managers, instead of the amount of available resources.
- Yarn containers may be killed when the used resources exceed the allocated ones.



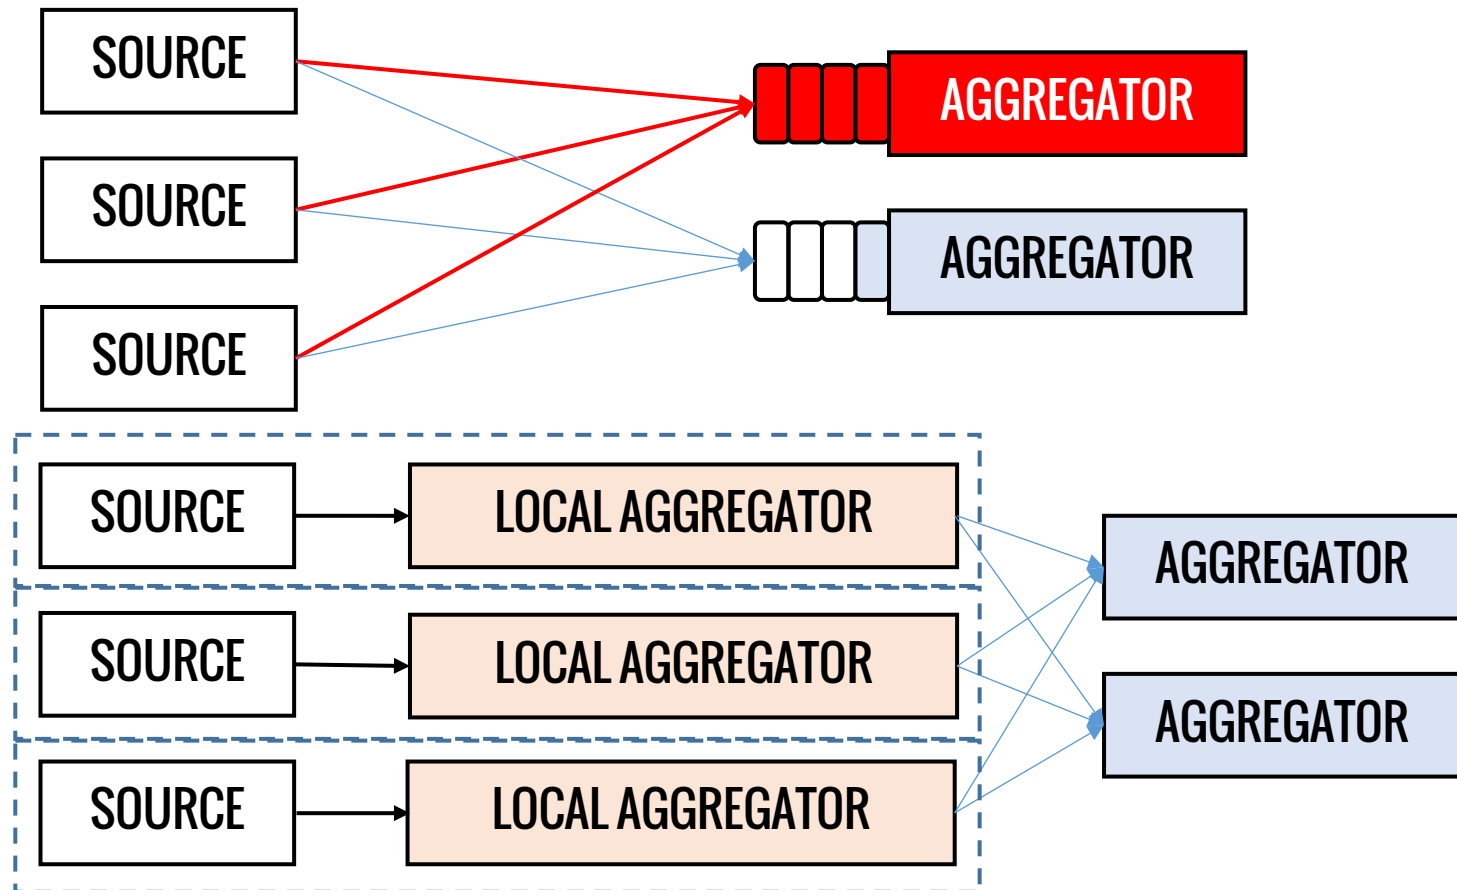
- Users can specify the resources needed by operators.
- A slot's resources are calculated by accumulating the resources of the operators in the slot.

- Slot instances are created and destroyed dynamically.
- A task manager creates a slot instance if its available resources are sufficient for the slot.
- A task manager destroys the slot instance if the tasks in the slot finish.

Local keyed streams

16




Current Problems: Performance is significantly degraded by data skew.



`words.keyBy().count()`

`words`
`.localKeyBy().window().count()`
`.keyBy().sum()`

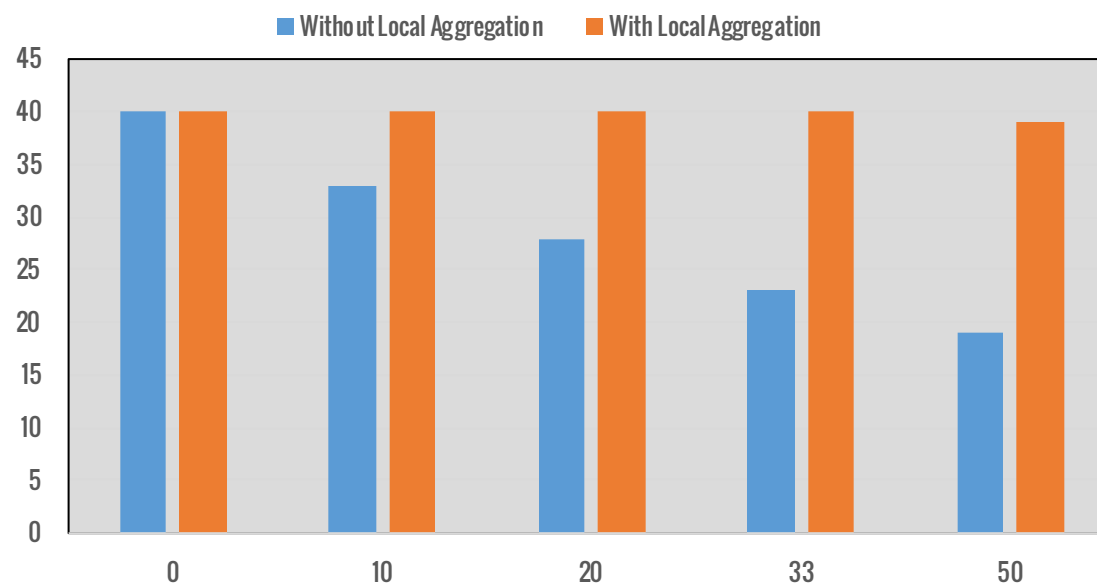
Local keyed streams



Each task has a complete key group range.

Groups are distributed to tasks according the number.

Groups with the same id are merged at restoring.



UDX

More than 40 UDX
are provided


Dim Join

Optimized implement
of joins with external
storage

Top N

Incremental Window

Allow uses to obtain partial
results of windows





Improve scheduling efficiency

Unified checkpoint mechanism for both streaming and batch jobs

Incorporating partitioning and timing into optimizer

SuperSQL: efficient data analytics across data sources (Hive, HBase, PostgreSQL, etc) and data centers

Thank You