

Realtime Store Visit Predictions at Scale


Luca Giovagnoli
luca@yelp.com



Outline



Product



Algorithm



Flink Implementation



Impact



Yelp's Mission

Connecting people with great local businesses.




12:40

82%

Fri, Mar 8



 Yelp • 18m ^

Luca, near Creator? 📍

Anna said, "My friends and I ordered the Creator vs. the World and Tumami ..." Read 140 other reviews here 📖

Manage notifications

Clear all



Phone



Hangouts



Games



Chrome



Clock




Photos





Creator

The wreck Room

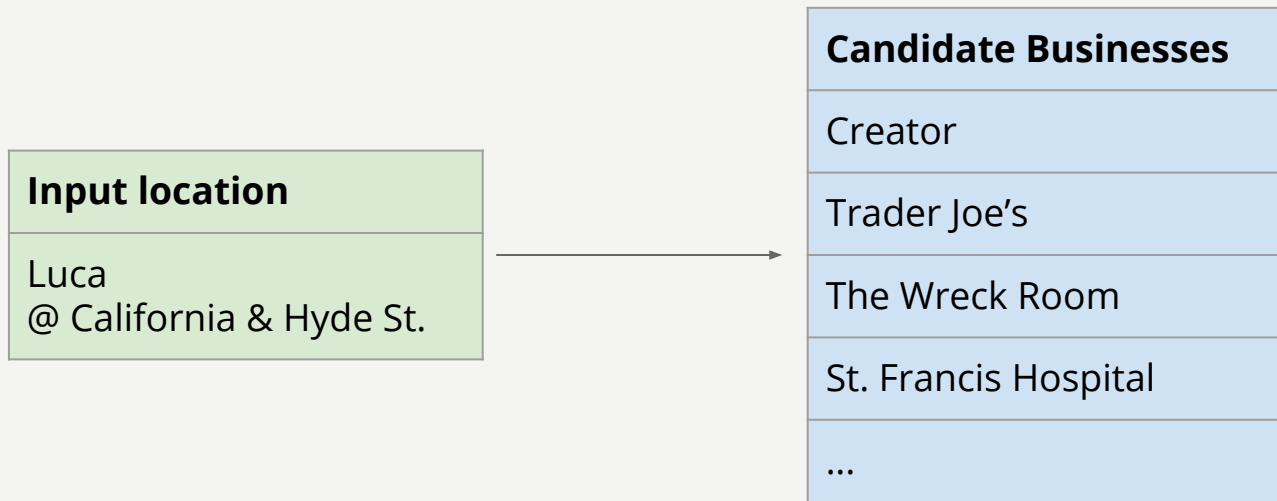


Trader Joe's

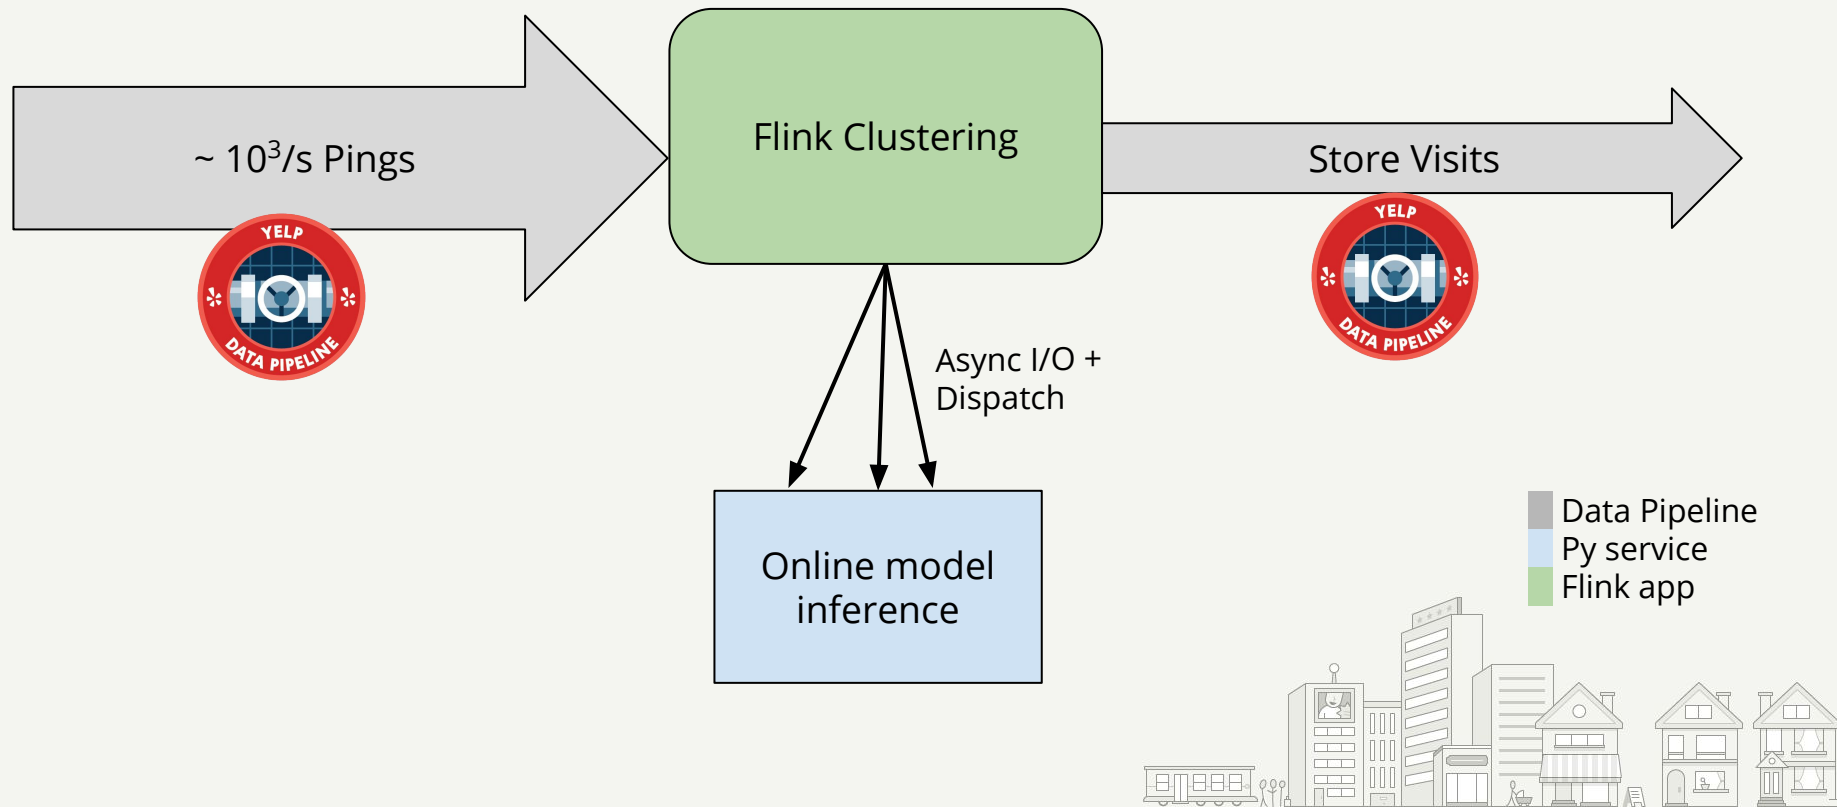
St Francis
Memorial Hospital




XGBoost ML model

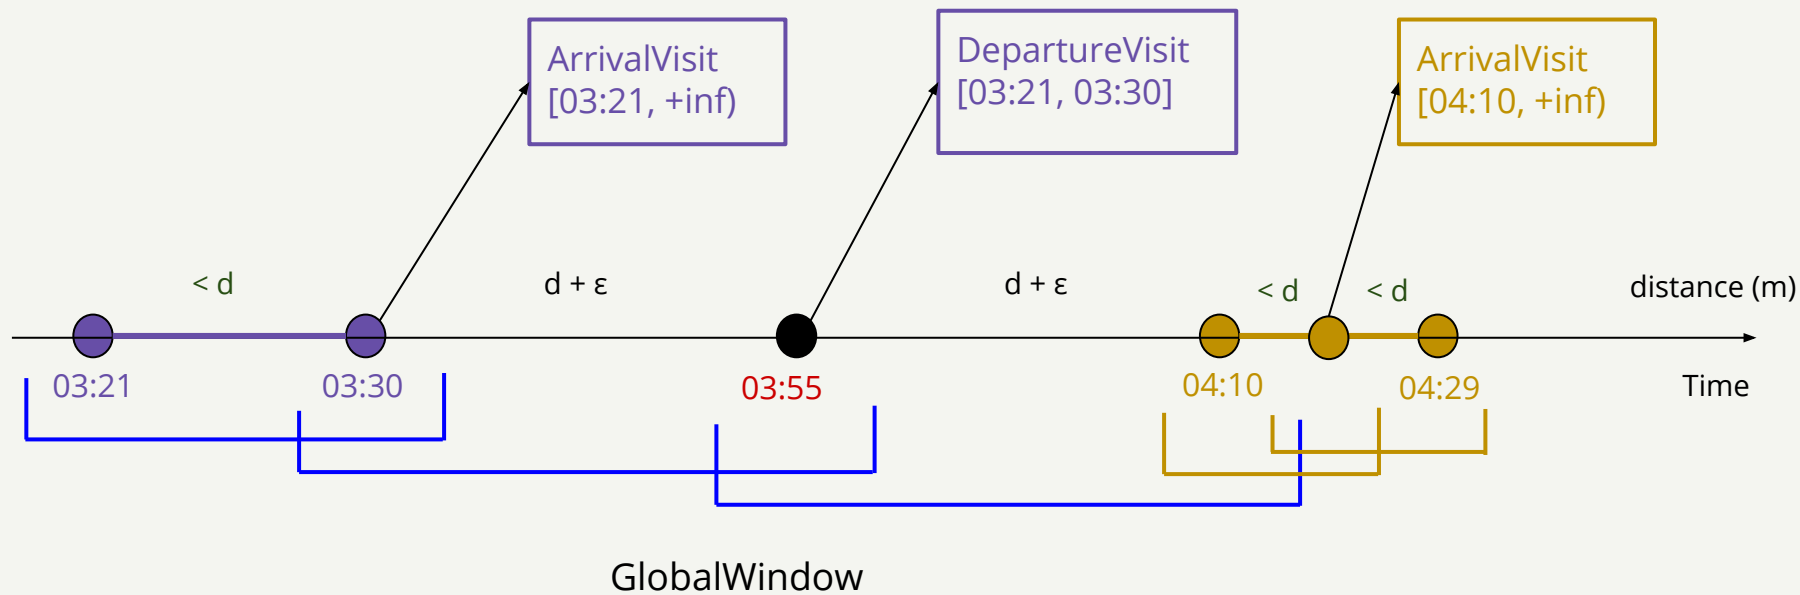


System overview




Clustering --> ProcessWindowFunction







Custom Trigger




Count Evictor

CountEvictor removes the oldest ping





- Mobiles go offline
- Mobiles clocks can be off



* Picture from <https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101> by Tyler Akidau





Location (Input)	Candidate businesses	Features			Confidence score
		Distance	Directions	...	
Luca @ California & Hyde St.	Creator	0.001	0.026	...	0.99
	Trader Joe's	0.021	0.071	...	0.13
	St. Francis Hospital



elasticsearch



Origins story



- Fragile in-house state
- No concurrency, no scaling
- At-least-once guarantee --> duplicate pushes



Impact



x10 Visits recall increase



up to $\sim 10^3$ ML predictions / sec



14 Flink instances down from $\sim 10^2$ Python



\sim **x5** times cheaper





We're Hiring!

www.yelp.com/careers/



@YelpEngineering



engineeringblog.yelp.com



github.com/yelp



yelp.com/careers



